

dataArtisans




Apache Flink® Training

System Overview



Apache Flink

A stream processor
with many applications



1 year of Flink - code



April 2014

April 2015

Stratosphere accepted as Apache Incubator Project

16 Apr 2014

We are happy to announce that Stratosphere has been accepted as a project for the [Apache Incubator](#). The [proposal](#) has been accepted by the Incubator PMC members earlier this week. The Apache Incubator is the first step in the process of giving a project to the [Apache Software Foundation](#). While under incubation, the project will move to the Apache infrastructure and adopt the community-driven development principles of the Apache Foundation. Projects can graduate from incubation to become top-level projects if they show activity, a healthy community dynamic, and releases.

We are glad to have Alan Gates as champion on board, as well as a set of great mentors, including Sean Owen, Ted Dunning, Owen O'Malley, Henry Saputra, and Ashutosh Chauhan. We are confident that we will make this a great open source effort.

0 Comments Apache Flink Login

Recommend Share Sort by Best

Start the discussion...

DataSet API (Java/Scala)

Flink core

Local

Remote

Yarn

Hadoop M/R

Python

Gelly

Table

ML

Dataflow

MRQL

Table

SAMOA

Dataflow

DataSet (Java/Scala)

DataStream (Java/Scala)

Flink core

Local


Remote

Yarn

Tez

Embedded

Flink Community



In top 5 of Apache's big data projects after one year in the Apache Software Foundation

The Apache Way




- Flink is an Apache **top-level project**




- Independent, **non-profit** organization
- **Community-driven** open source software development approach
- **Public communication** and open to new contributors

What is Apache Flink?



DataSet (Java/Scala/Python)

DataStream (Java/Scala)



Local

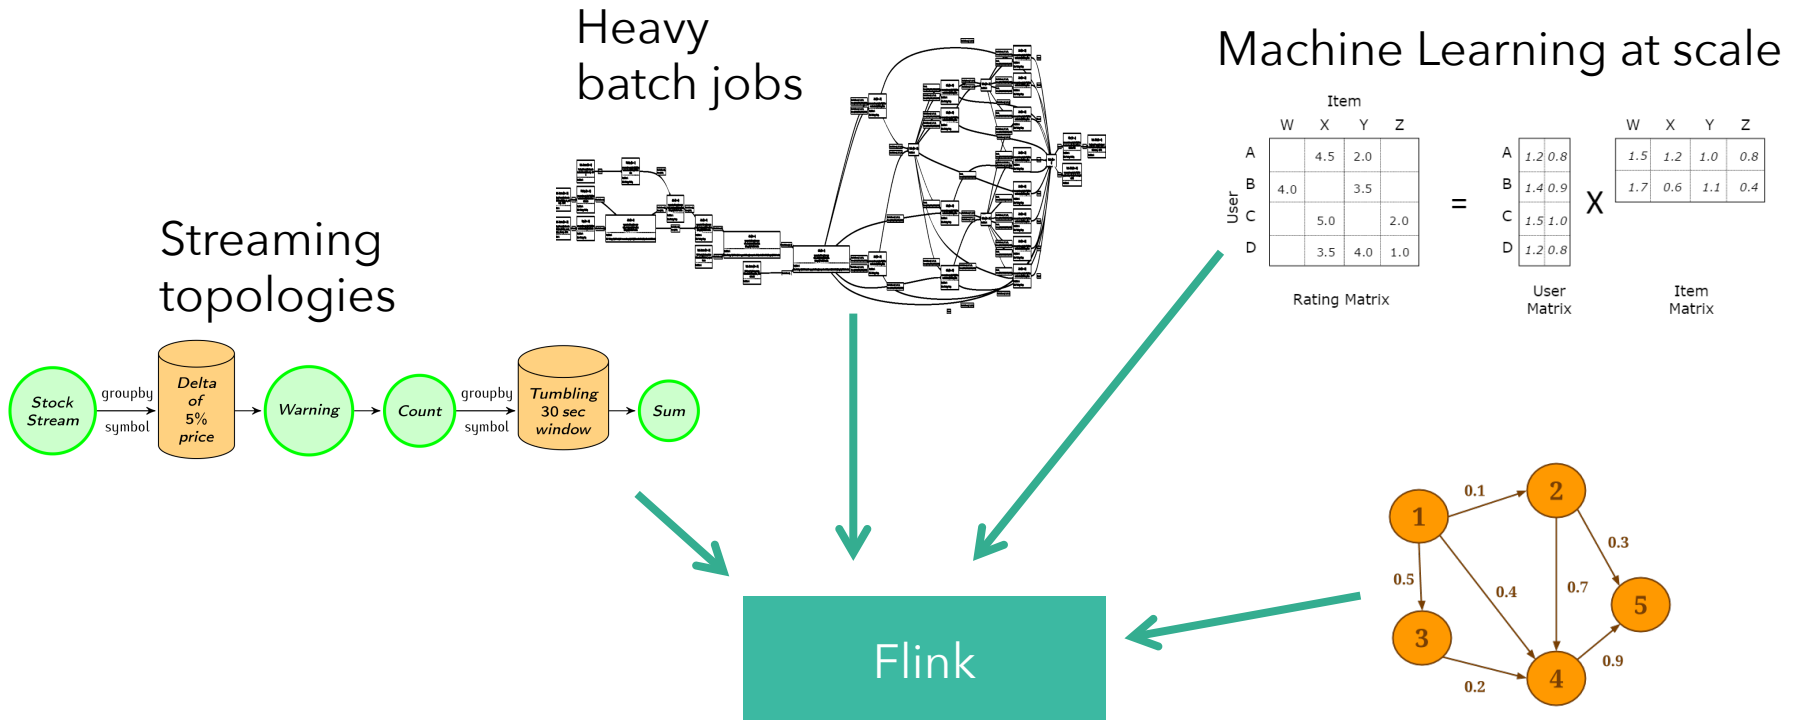
Remote

Yarn

Tez

Embedded

Native workload support




How can an engine **natively** support all these workloads?

And what does native **mean**?

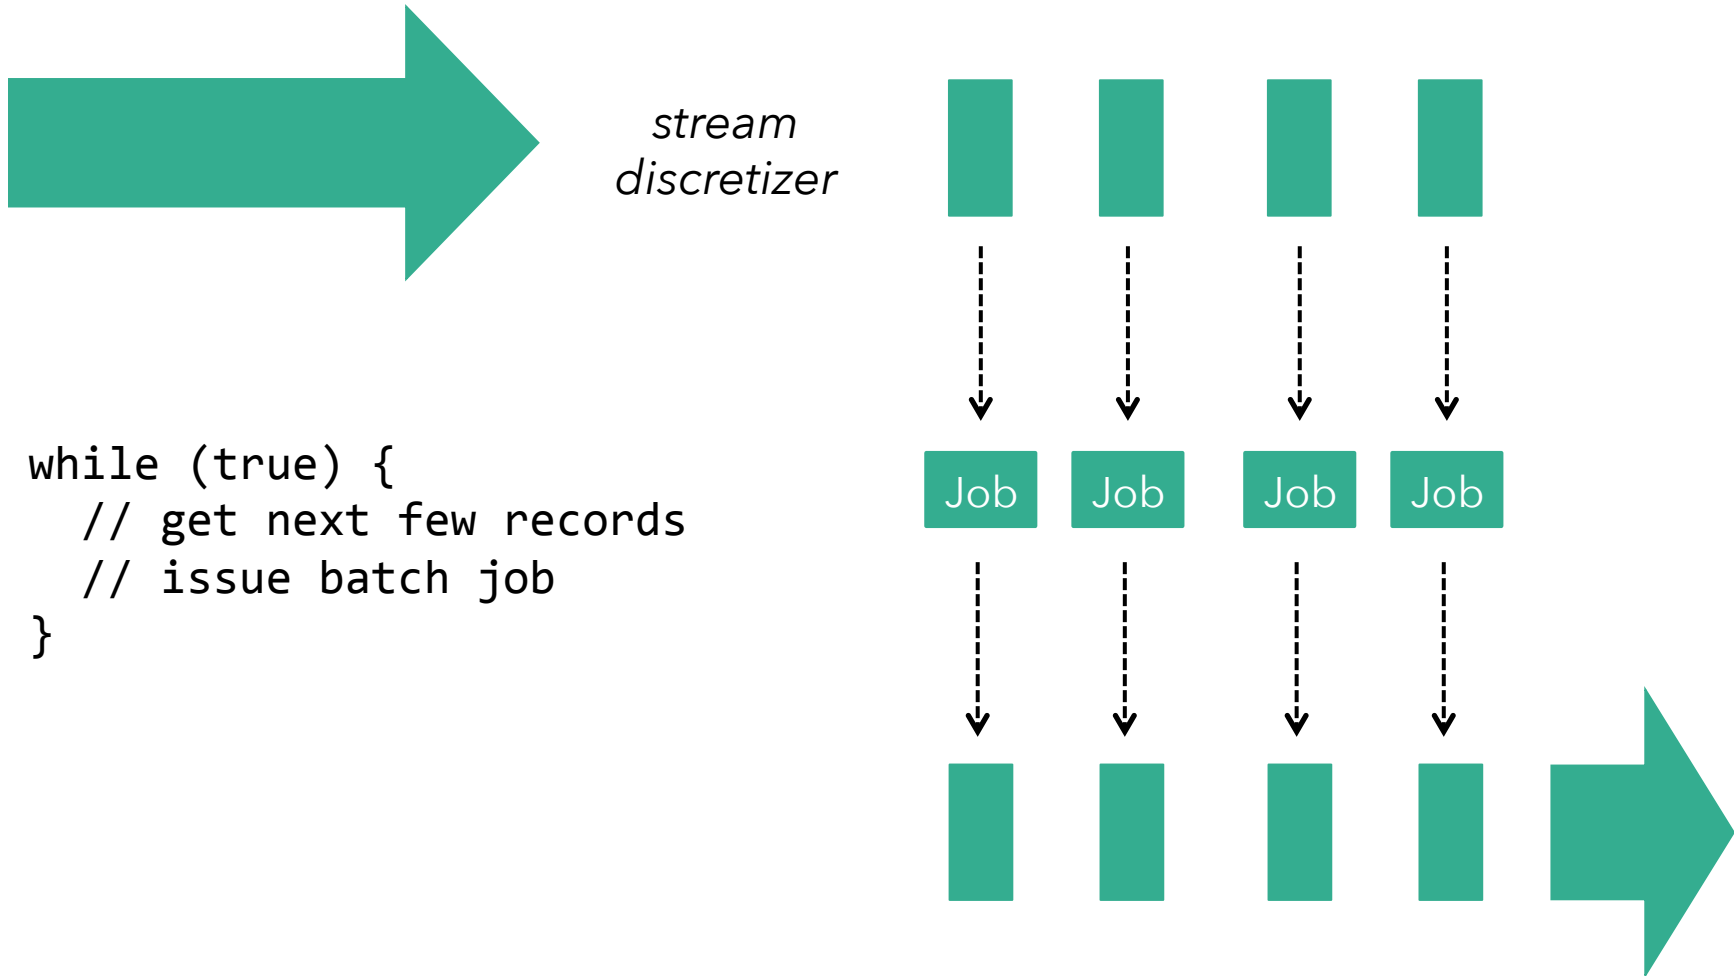
E.g.: Non-native iterations



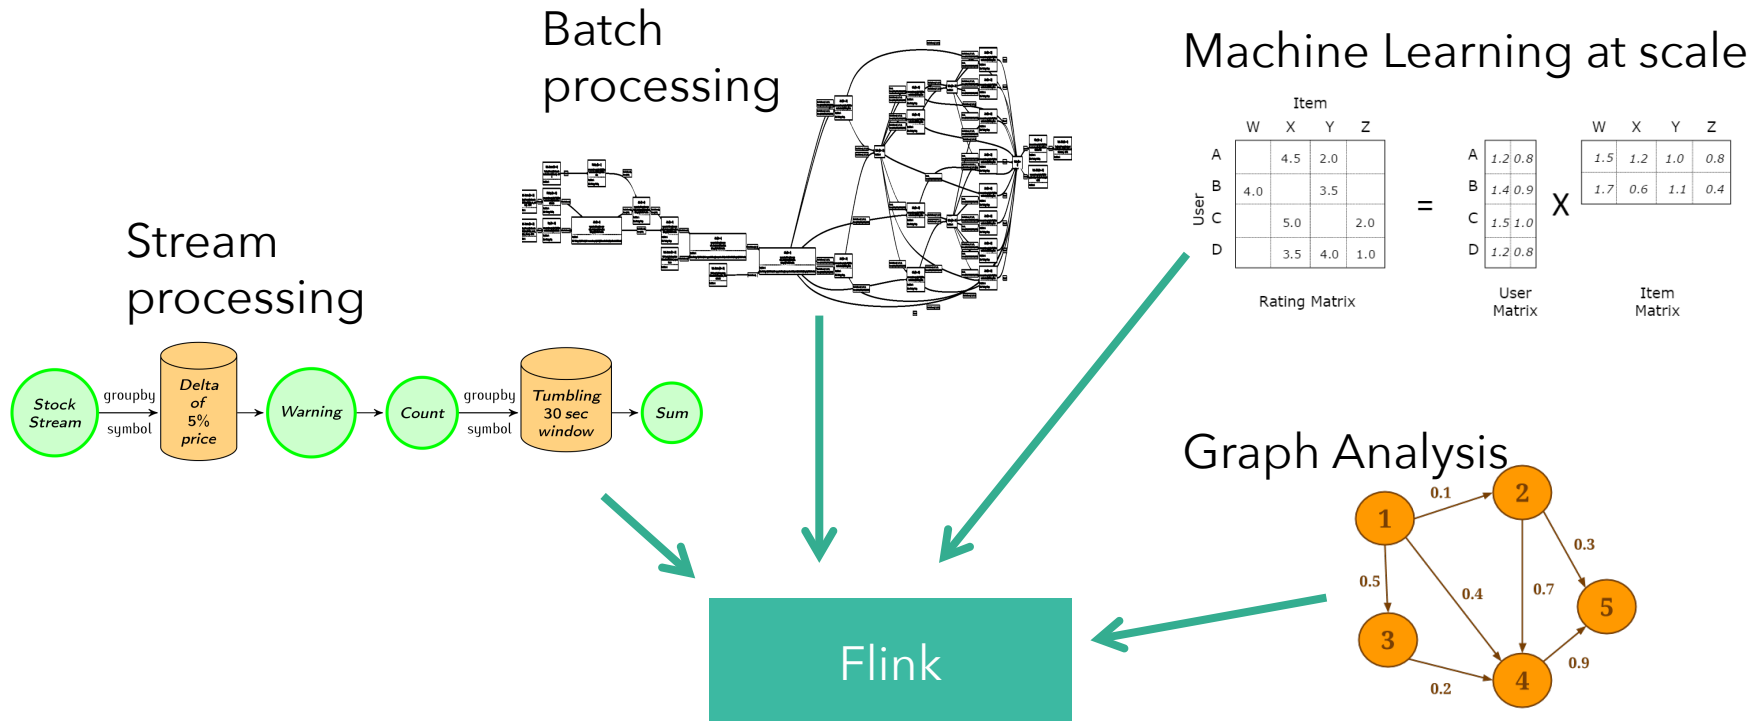
```
for (int i = 0; i < maxIterations; i++) {  
    // Execute MapReduce job  
}
```



E.g.: Non-native streaming



Native workload support








How can an engine **natively** support all these workloads?
And what does "native" **mean**?

Flink Engine

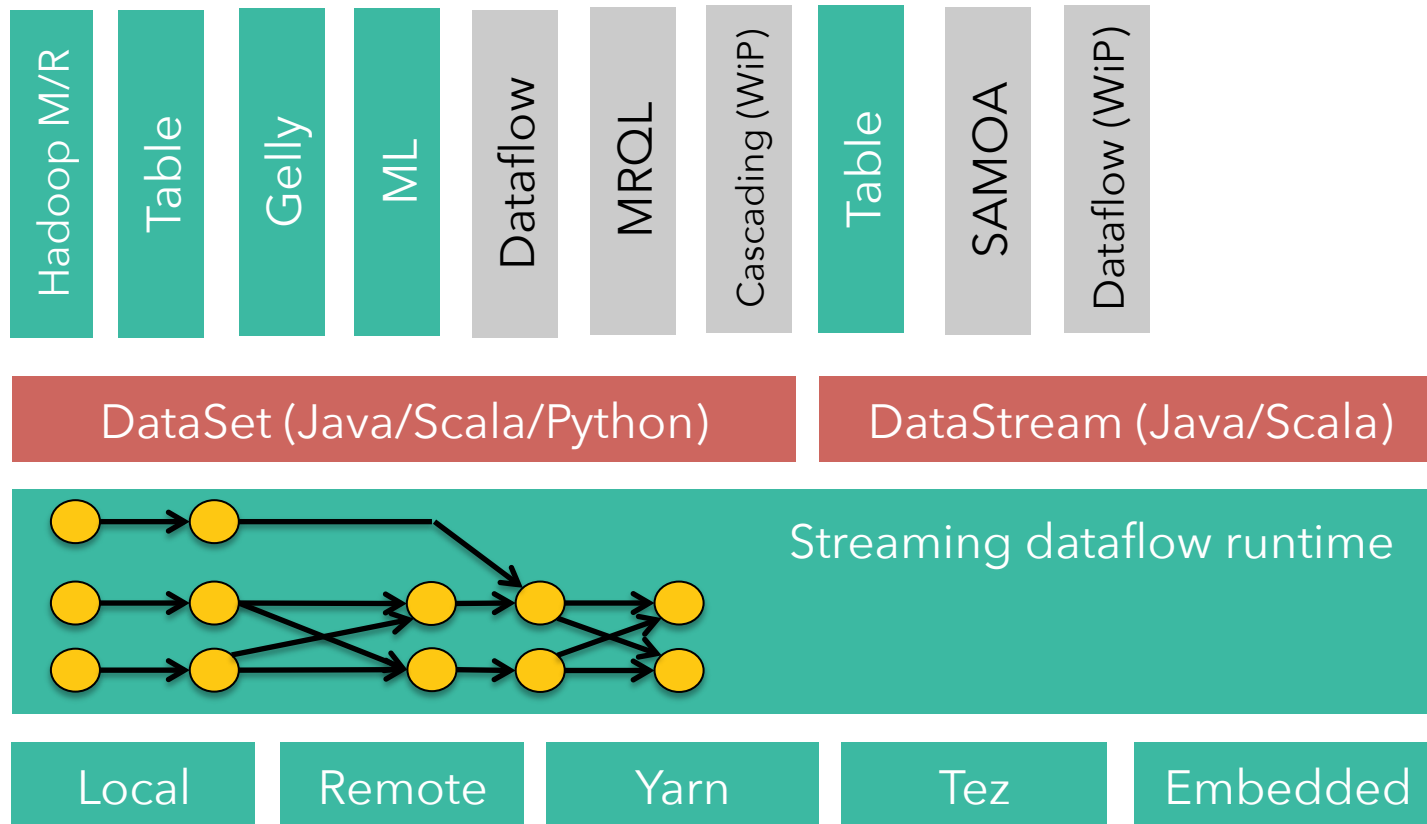


1. Execute everything as streams
2. Iterative (cyclic) dataflows
3. Mutable state
4. Operate on managed memory
5. Special code paths for batch




What is a Flink Program?

Flink stack



Basic API Concept



How do I write a Flink program?


1. Bootstrap sources
2. Apply operations
3. Output to source

Batch & Stream Processing



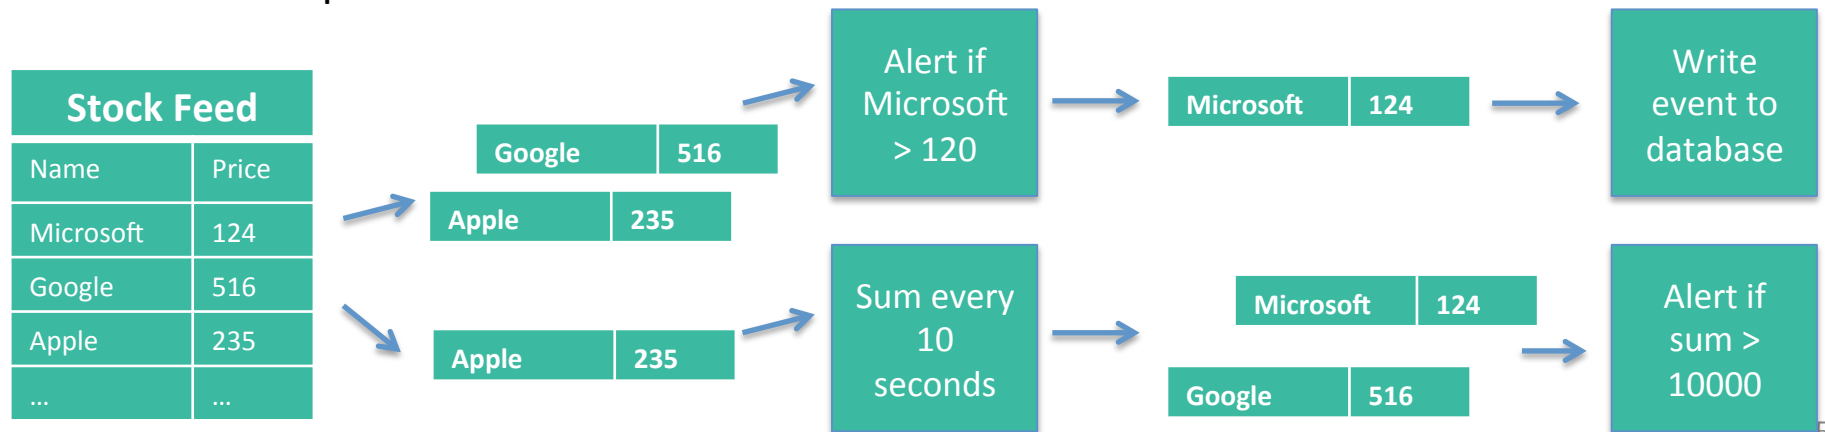
- DataSet API

Example: Map/Reduce paradigm



- DataStream API

Example: Live Stock Feed





Streaming & Batch

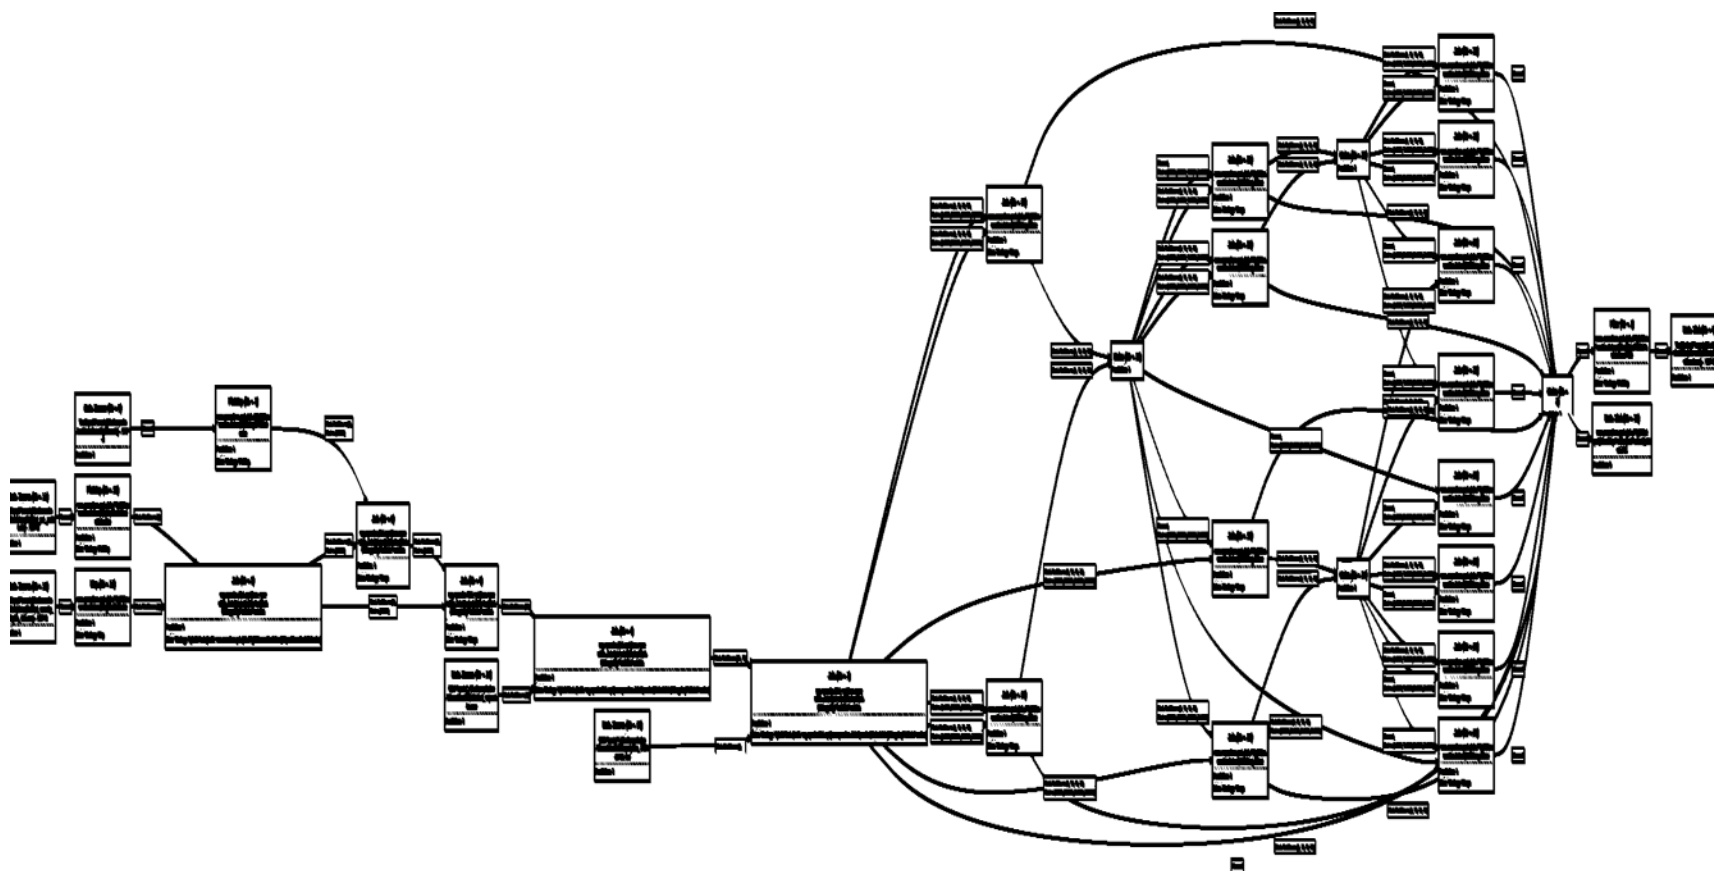


	Streaming	Batch
Input	infinite	finite
Data transfer	pipelined	blocking or pipelined
Latency	low	high

Scaling out



Scaling up



Sources (selection)



Collection-based

- fromCollection
- fromElements

File-based

- TextInputFormat
- CsvInputFormat

Other

- SocketInputFormat
- KafkaInputFormat
- Databases

Sinks (selection)



File-based

- TextOutputFormat
- CsvOutputFormat
- PrintOutput

Others

- SocketOutputFormat
- KafkaOutputFormat
- Databases

Hadoop Integration



Out of the box

- Access HDFS
- Yarn Execution (covered later)
- Reuse data types (Writables)

With a thin wrapper

- Reuse Hadoop input and output formats
- Reuse functions like Map and Reduce


What's the Lifecycle of a Program?

From Program to Dataflow




```
case class Path (from: Long, to: Long)
val tc = edges.iterate(10) {
  paths: DataSet[Path] =>
    val next = paths
      .join(edges)
      .where("to")
      .equalTo("from") {
        (path, edge) =>
          Path(path.from, edge.to)
      }
      .union(paths)
      .distinct()
    next
}
```



Program



Pre-flight (Client)




*Dataflow
Graph*



Master

*deploy
operators*

*track
intermediate
results*



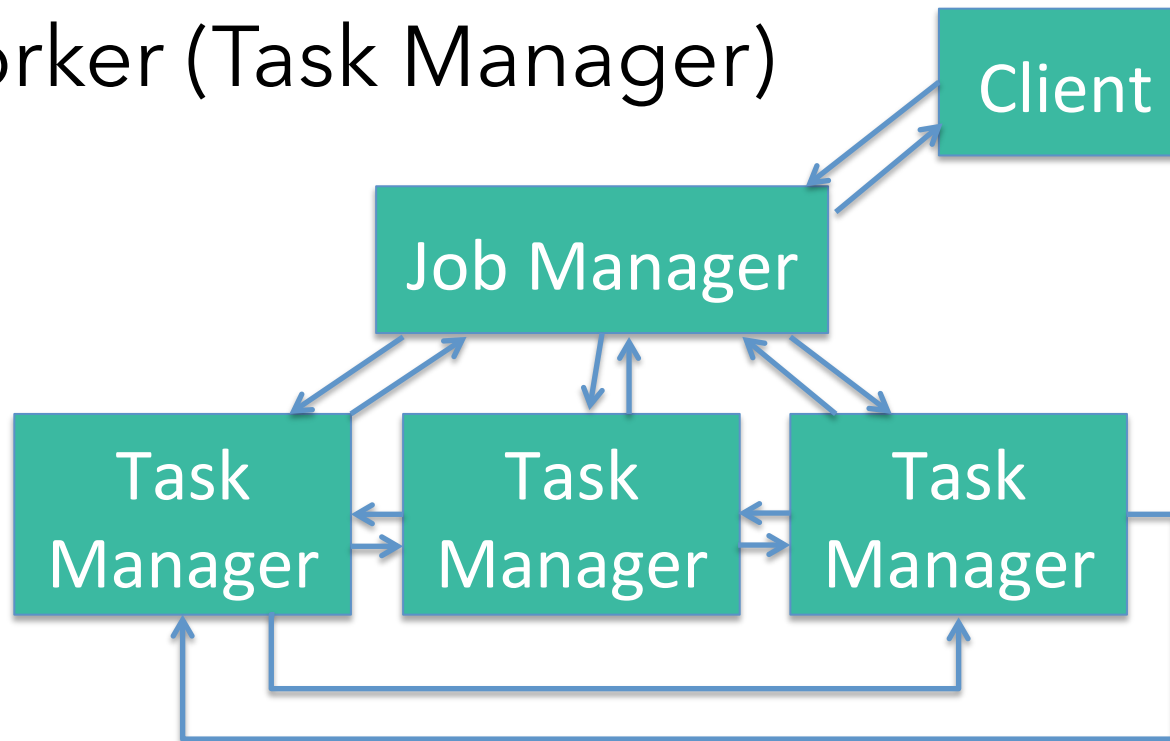
Workers

23

Architecture Overview



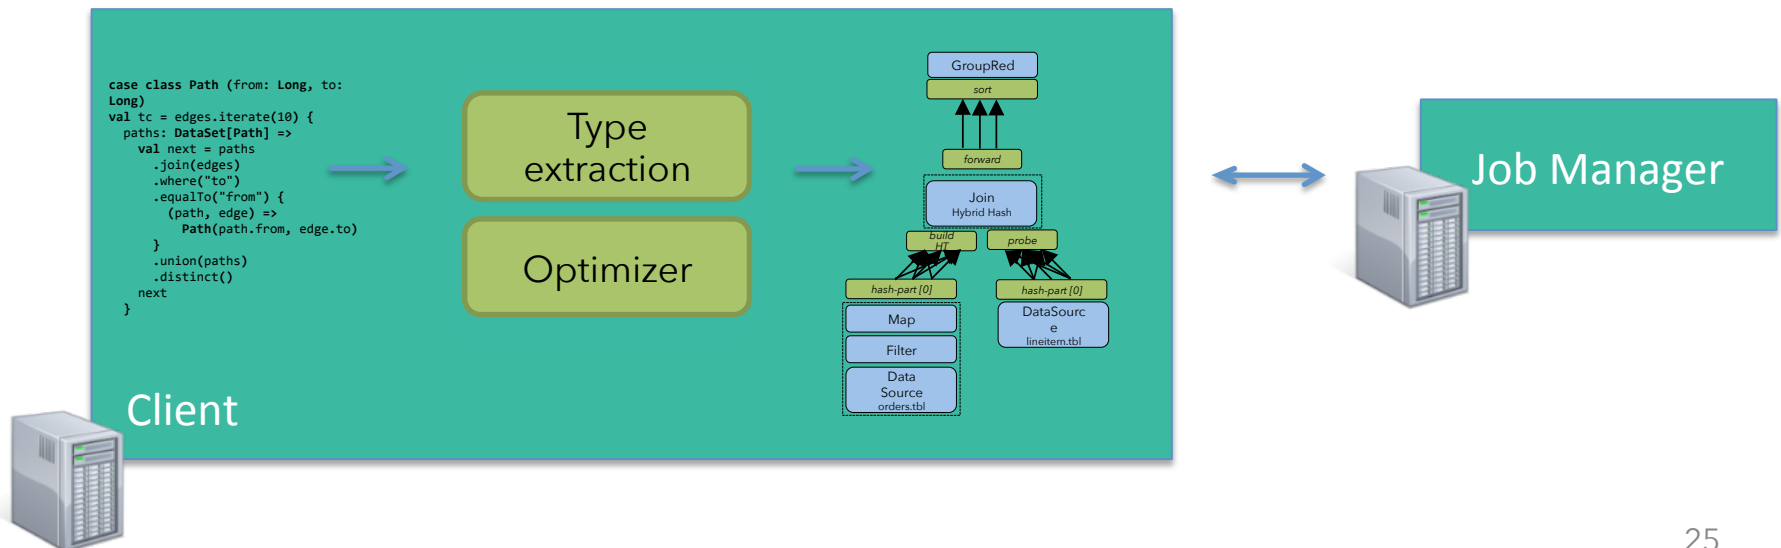
- Client
- Master (Job Manager)
- Worker (Task Manager)



Client



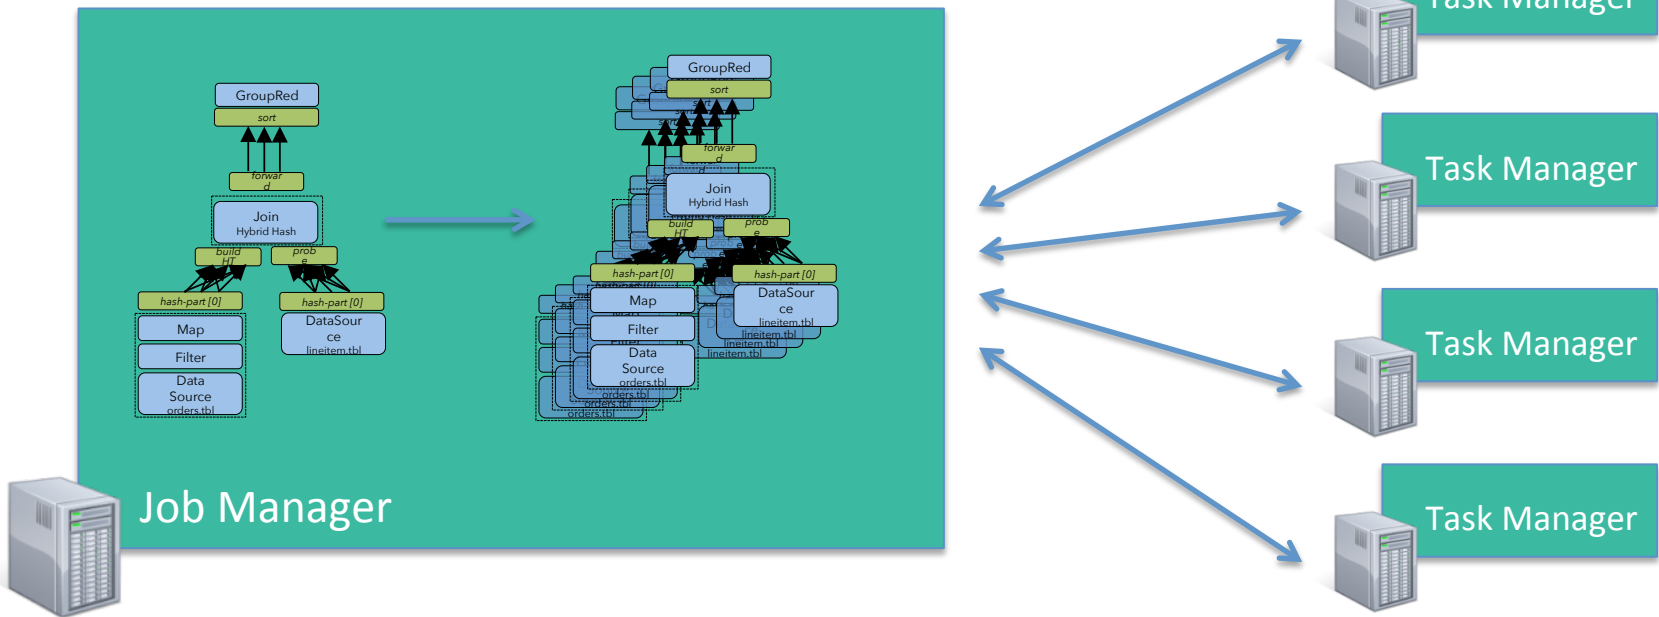
- Optimize
- Construct job graph
- Pass job graph to job manager
- Retrieve job results



Job Manager




- **Parallelization:** Create Execution Graph
- **Scheduling:** Assign tasks to task managers
- **State:** Supervise the execution



Task Manager

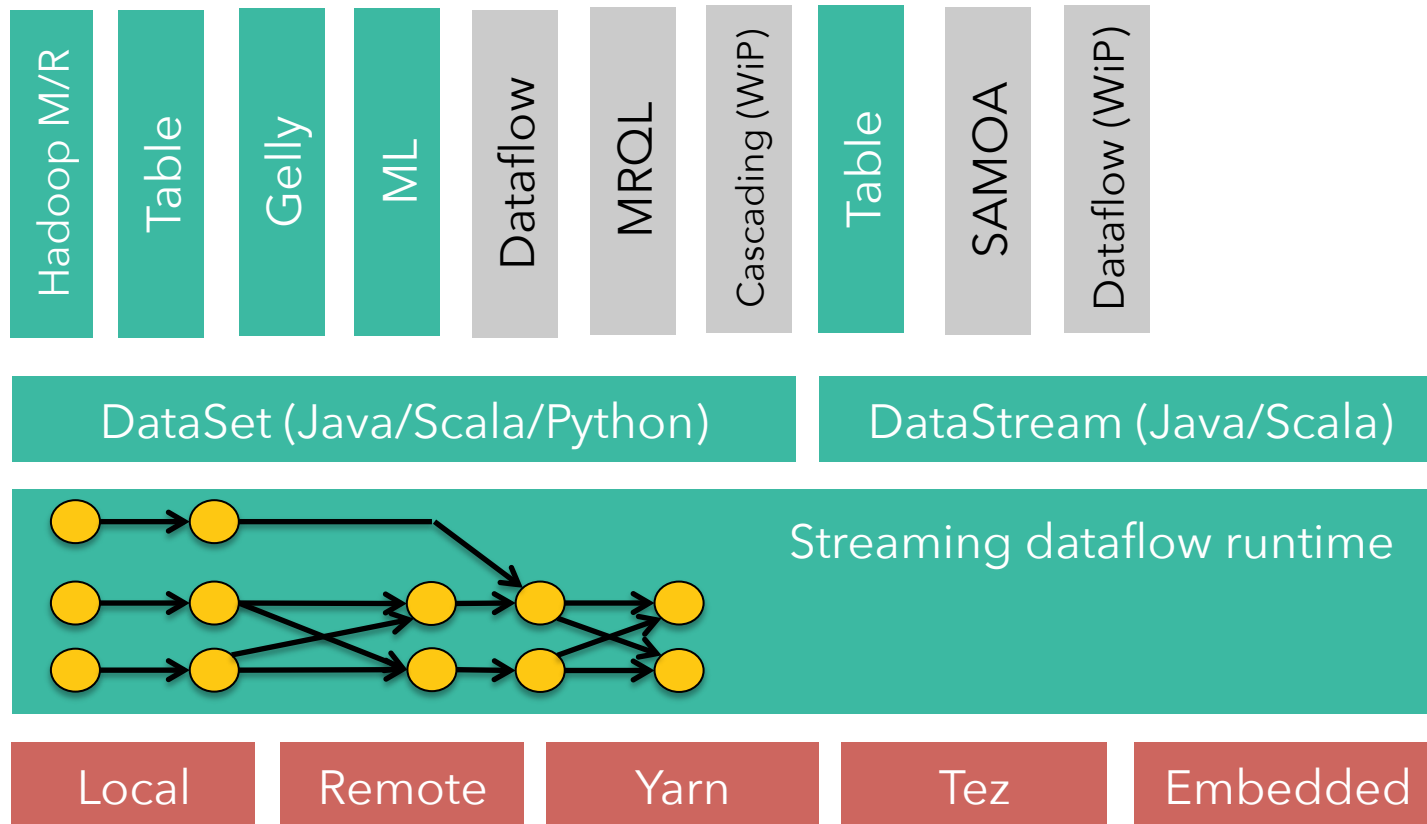


- Operations are split up into **tasks** depending on the specified parallelism
- Each parallel instance of an operation runs in a separate **task slot**
- The scheduler may run several tasks from different operators in one task slot



Execution Setups

Ways to Run a Flink Program



Local Execution



- Starts local Flink cluster
- All processes run in the same JVM
- Behaves just like a regular Cluster
- Very useful for developing and debugging



Embedded Execution

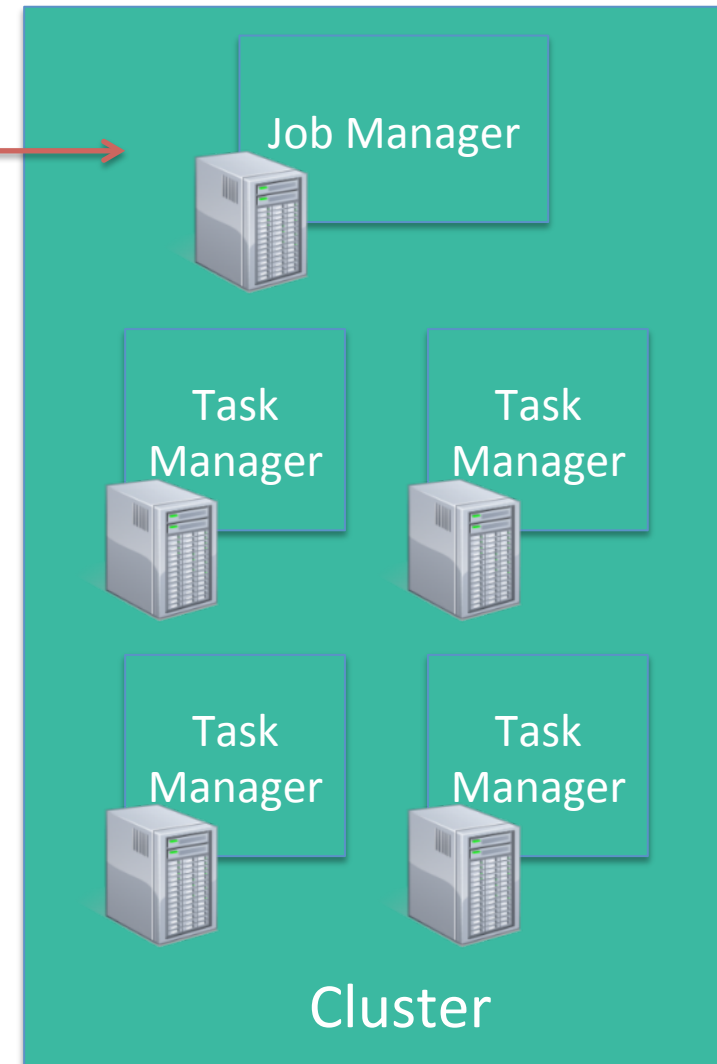


- Runs operators on simple **Java collections**
- Lower overhead
- Does not use memory management
- Useful for testing and debugging

Remote Execution



Submit job

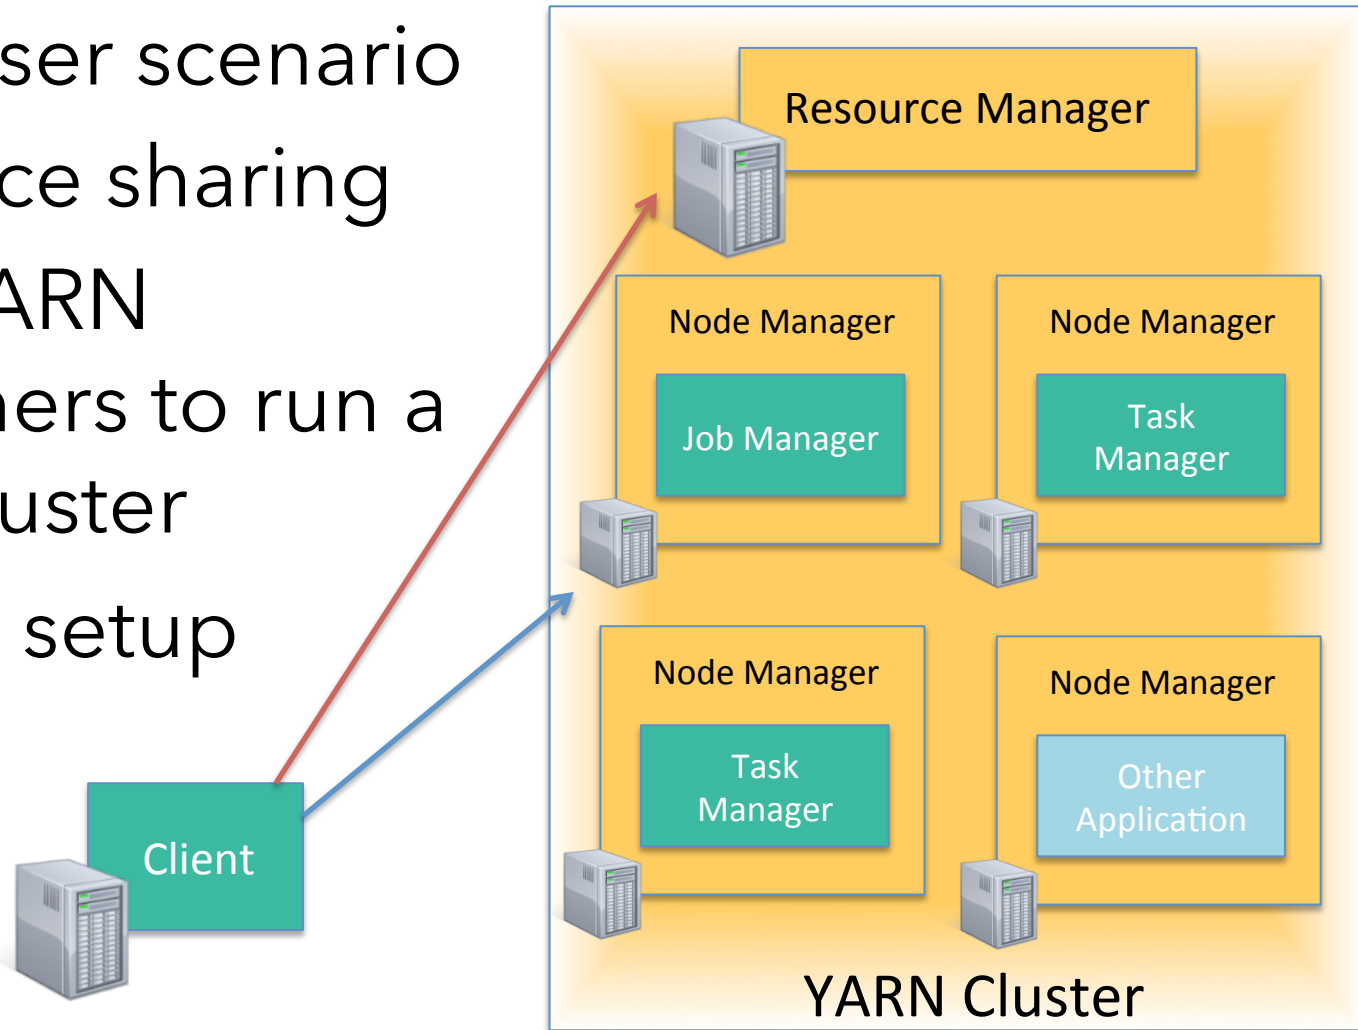


- **Submit** a Job remotely
- **Monitor** the status of a job

YARN Execution



- Multi-user scenario
- Resource sharing
- Uses YARN containers to run a Flink cluster
- Easy to setup





Execution




- Leverages Apache Tez's runtime
- Built on top of YARN
- Good YARN citizen
- Fast path to elastic deployments
- Slower than native Flink

Flink compared to other projects

Batch & Streaming projects



Batch only




Streaming only



Hybrid



Batch comparison



API	low-level	high-level	high-level
Data Transfer	batch	batch	pipelined & batch
Memory Management	disk-based	JVM-managed	Active managed
Iterations	file system cached	in-memory cached	streamed
Fault tolerance	task level	task level	job level
Good at	massive scale out	data exploration	heavy backend & iterative jobs
Libraries	many external	built-in & external	evolving built-in & external

Streaming comparison



Streaming	“true”	mini batches	“true”
API	low-level	high-level	high-level
Fault tolerance	tuple-level ACKs	RDD-based (lineage)	coarse checkpointing
State	not built-in	external	internal
Exactly once	at least once	exactly once	exactly once
Windowing	not built-in	restricted	flexible
Latency	low	medium	low
Throughput	medium	high	high

Thank you for listening!