

dataArtisans



Apache Flink® Training

DataStream API Advanced

December 10, 2015

What kind of data can Flink handle?

Type System and Keys

Apache Flink's Type System



- Flink aims to support all data types
 - Ease of programming
 - Seamless integration with existing code
 - DataSet and DataStream API share the same type system!
- Programs are analyzed before execution
 - Used data types are identified
 - Serializer & comparator are configured

Apache Flink's Type System



- Data types are either
 - Atomic types (like Java Primitives)
 - Composite types (like Flink Tuples)
- Composite types nest other types
- Not all data types can be used as keys!
 - Flink partitions DataStreams on keys
 - Key types must be comparable

Atomic Types



| Flink Type | Java Type | Can be used as key? |
|--------------|---|--|
| BasicType | Java Primitives (Integer, String, ...) | Yes |
| ArrayType | Arrays of Java primitives or objects | Primitive Arrays: Yes Object Arrays: No |
| WritableType | Implements Hadoop's Writable interface | Yes, if implements WritableComparable |
| GenericType | Any other type | Yes, if implements Comparable |

Composite Types



- Are composed of fields with other types
 - Fields types can be atomic or composite
- Fields can be addressed as keys
 - Field type must be a key type!
- A composite type can be a key type
 - All field types must be key types!

TupleType



- Java:
`org.apache.flink.api.java.tuple.Tuple1` to `Tuple25`
- Scala:
use default Scala tuples (1 to 22 fields)
- Tuple fields are typed

```
Tuple3<Integer, String, Double> t3 =  
    new Tuple3<>(1, "2", 3.0);
```

```
val t3: (Int, String, Double) = (1, "2", 3.0)
```

- Tuples give the best performance

TupleType



- Define keys by field position

```
DataStream<Tuple3<Integer, String, Double>> d = ...  
// key stream by String field  
d.keyBy(1);
```

- Or field names

```
// key stream by Double field  
d.keyBy("f2");
```


PojoType



- Any Java class that
 - Has an empty default constructor
 - Has publicly accessible fields
(public field or default getter & setter)

```
public class Person {  
    public int id;  
    public String name;  
    public Person() {};  
    public Person(int id, String name) {...};  
}
```

```
DataStream<Person> p =  
    env.fromElements(new Person(1, "Bob"));
```

PojoType



- Define keys by field name

```
DataStream<Person> p = ...  
// key stream by "name" field  
p.keyBy("name");
```

Scala CaseClasses



- Scala case classes are natively supported

```
case class Person(id: Int, name: String)
d: DataStream[Person] =
    env.fromElements(Person(1, "Bob"))
```

- Define keys by field name

```
// key stream by field "name"
d.keyBy("name")
```

Composite & Nested Keys



```
DataStream<Tuple3<String, Person, Double>> d;
```

- Composite keys are supported

```
// key stream by both long fields  
d.keyBy(0, 1);
```

- Nested fields can be used as types

```
// key stream by nested "name" field  
d.keyBy("f1.name");
```

- Full types can be used as key using "*" wildcard

```
// key stream by complete nested Pojo field  
d.keyBy("f1.*");
```

- "*" wildcard can also be used for atomic types

KeySelectors



- Keys can be computed using KeySelectors

```
public class SumKeySelector implements
    KeySelector
```

```
DataStream<Tuple2<Long, Long>> d = ...
d.keyBy(new SumKeySelector());
```

Windows and Aggregates

Windows



- Aggregations on DataStreams are different from aggregations on DataSets
 - e.g., it is not possible to count all elements of an unbounded DataStream
- DataStream aggregations make sense on windowed streams
 - i.e., a finite set of stream elements
- Only windows on keyed stream can be processed in parallel

Windows (2)



```
// (age, count) of passengers  
DataStream<Tuple2<Integer, Integer>> passengers = ...
```

Passengers

```
// group by first field (age)  
    .keyBy(0)  
// window of 1 minute length triggered every 10 seconds  
    .timeWindow(Time.minutes(1), Time.seconds(10))  
// sum values of second field (count)  
    .sum(2);
```


Predefined Keyed Windows



- Tumbling time window
`.timeWindow(Time.minutes(1))`
- Sliding time window
`.timeWindow(Time.minutes(1), Time.seconds(30))`
- Tumbling count window
`.countWindow(100)`
- Sliding count window
`.countWindow(100, 10)`

Aggregations on Windowed Streams



```
// (name, age) of passengers
DataStream<Tuple2<String, Integer>> passengers = ...

passengers
    // group by first field (age)
    .keyBy(0)
    // window of 1 minute length triggered every 10 seconds
    .timeWindow(Time.minutes(1), Time.seconds(10))
    // apply a custom window function on window data
    .apply(new CountByAge());
```

MapWindow



```
public static class CountByAge implements WindowFunction<
    Tuple2<String, Integer>, // input type
    Tuple3<Integer, Long, Integer>, // output type
    Tuple, // key type
    TimeWindow> // window type
{
    @Override
    public void apply(
        Tuple key,
        TimeWindow window,
        Iterable<Tuple2<String, Integer>> persons,
        Collector<Tuple3<Integer, Long, Integer>> out) {

        int age = ((Tuple1<Integer>)key).f0;
        int cnt = 0;

        for (Tuple2<String, Integer> p : persons) {
            cnt++;
        }
        // return (age, window-end-time, count)
        out.collect(new Tuple3<>(age, window.getEnd(), cnt));
    }
}
```

Operations on Windowed Streams



- `reduce(reduceFunction)`
 - Apply a functional reduce function to the window
- `fold(inialVal, foldFunction)`
 - Apply a functional fold function with a specified initial value to the window
- Aggregation functions
 - `sum()`, `min()`, `max()`, and others

Windows on non-keyed streams



- Windows on non-keyed streams are not parallel!
- TimeWindow (tumbling, 10 seconds)
 `.timeWindowAll(Time.seconds(10))`
- CountWindow (sliding, 20/10)
 `.countWindowAll(20, 10)`

Custom window logic



- The DataStream API allows to define very custom window logic
- Trigger
 - defines when to evaluate a window
 - whether to purge the window or not
- Evictor
 - Allows to remove elements from a window before it is evaluated
- *Careful!* This part of the API requires a good understanding of the windowing mechanism!

Working With Multiple Streams

Connecting Streams



- *Connect* two DataStreams to correlated them with each other
- Apply functions on connected streams to share state

```
DataStream<String> strings = ...
```

```
DataStream<Integer> ints = ...
```

```
ConnectedStreams<String, Integer> coStream =  
    strings.connect(ints);
```


Map on Connected Streams



```
DataStream<String> strings = ...
```

```
DataStream<Integer> ints = ...
```

```
ints.connect(strings)
```

```
    .map(new CoMapFunction<Integer, String, Boolean>() {
```

```
        @Override
```

```
        public Boolean map1 (Integer value) {
```

```
            return true;
```

```
        }
```

```
        @Override
```

```
        public Boolean map2 (String value) {
```

```
            return false;
```

```
        }
```

```
    });
```

FlatMap on Connected Streams



```
DataStream<String> strings = ...
```

```
DataStream<Integer> ints = ...
```

```
ints.connect(strings)
```

```
    .flatMap(new CoFlatMapFunction<Integer,String,String>() {
```

```
        @Override
```

```
        public void flatMap1 (Integer value, Collector<String> out) {  
            out.collect(value.toString());
```

```
        }
```

```
        @Override
```

```
        public void flatMap2 (String value, Collector<String> out) {  
            for (String word: value.split(" ")) {  
                out.collect(word);
```

```
            }
```

```
        }
```

```
    });
```

Connecting to Apache Kafka

Kafka and Flink



- “Apache Kafka is a distributed, partitioned, replicated commit log service”
- Kafka uses Apache Zookeeper for coordination
- Kafka maintains feeds of messages in categories called topics
- A Kafka topic can be read by Flink to produce a DataStream, and a DataStream can be written to a Kafka topic
- Flink coordinates with Kafka to provide recovery in the case of failures

Reading Data from Kafka



- Enable checkpointing
E.g., `env.enableCheckpointing(5000);`
- Add a `DataStream` source from a Kafka topic

```
Properties props = new Properties();
props.setProperty("zookeeper.connect", "localhost:2181");
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("group.id", "myGroup");

// create a data source
DataStream<MyData> rides = env.addSource(
    new FlinkKafkaConsumer082<MyData>(
        "myTopic",
        new MyDataSchema(),
        props)
);
```

Writing Data to Kafka



- Add a Kafka sink to a DataStream by providing
 - The broker address
 - The topic name
 - A serialization schema

```
DataStream<String> aStream = ...  
aStream.addSink(  
    new FlinkKafkaProducer<String>(  
        "localhost:9092", // default local broker  
        "myTopic",  
        new SimpleStringSchema));
```

More API Features

Not Covered Here



- Iterations (feedback edges)
 - Enables certain machine learning algorithms
- More transformations
 - split, union, window join, ...

