data Artisans

Apache Flink® Training

DataStream API Basic

December 10, 2015

DataStream API

- Stream Processing

- Java and Scala

- All examples here in Java for Flink 1.0

- Documentation available at
  flink.apache.org

DataStream API by Example

Window WordCount: main Method

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Stream Execution Environment

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Data Sources

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Data types

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Transformations

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

User functions

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

DataSinks

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Execute!

```java
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();
    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
            // read stream of words from socket
            .socketTextStream("localhost", 9999)
            // split up the lines in tuples containing: (word,1)
            .flatMap(new Splitter())
            // key stream by the tuple field "0"
            .keyBy(0)
            // compute counts every 5 minutes
            .timeWindow(Time.minutes(5))
            //sum up tuple field "1"
            .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Window WordCount: FlatMap

```java
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

    @Override
    public void flatMap(String value,
                        Collector<Tuple2<String, Integer>> out)
        throws Exception {
        // normalize and split the line
        String[] tokens = value.toLowerCase().split("\\W+");

        // emit the pairs
        for (String token : tokens) {
            if (token.length() > 0) {
                out.collect(
                    new Tuple2<String, Integer>(token, 1));
            }
        }
    }
}
```

WordCount: Map: Interface

```java
public static class Splitter
  implements FlatMapFunction<String, Tuple2<String, Integer>> {

  @Override
  public void flatMap(String value,
                      Collector<Tuple2<String, Integer>> out)
    throws Exception {
      // normalize and split the line
      String[] tokens = value.toLowerCase().split("\\W+");

      // emit the pairs
      for (String token : tokens) {
        if (token.length() > 0) {
          out.collect(
            new Tuple2<String, Integer>(token, 1));
        }
      }
    }
}
```

WordCount: Map: Types

```java
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

  @Override
  public void flatMap(String value,
                      Collector<Tuple2<String, Integer>> out)
    throws Exception {
      // normalize and split the line
      String[] tokens = value.toLowerCase().split("\\W+");

      // emit the pairs
      for (String token : tokens) {
        if (token.length() > 0) {
          out.collect(
            new Tuple2<String, Integer>(token, 1));
        }
      }
  }
}
```

WordCount: Map: Collector

```java
public static class Splitter
  implements FlatMapFunction<String, Tuple2<String, Integer>> {

  @Override
  public void flatMap(String value,
                    Collector<Tuple2<String, Integer>> out)
    throws Exception {
      // normalize and split the line
      String[] tokens = value.toLowerCase().split("\\W+");

      // emit the pairs
      for (String token : tokens) {
          if (token.length() > 0) {
              out.collect(
                  new Tuple2<String, Integer>(token, 1));
          }
      }
  }
}
```

DataStream API Concepts

(Selected) Data Types

- Basic Java Types
  - String, Long, Integer, Boolean,…
  - Arrays

- Composite Types
  - Tuples
  - Many more (covered in the advanced slides)

Tuples

- The easiest and most lightweight way of encapsulating data in Flink

- Tuple1 up to Tuple25

```
Tuple2<String, String> person = new Tuple2<>("Max", "Mustermann");

Tuple3<String, String, Integer> person = new Tuple3<>("Max", "Mustermann", 42);

Tuple4<String, String, Integer, Boolean> person =
  new Tuple4<>("Max", "Mustermann", 42, true);

// zero based index!
String firstName = person.f0;
String secondName = person.f1;
Integer age = person.f2;
Boolean fired = person.f3;
```

Transformations: Map

```java
DataStream<Integer> integers = env.fromElements(1, 2, 3, 4);

// Regular Map – Takes one element and produces one element
DataStream<Integer> doubleIntegers =
    integers.map(new MapFunction<Integer, Integer>() {
    @Override
        public Integer map(Integer value) {
            return value * 2;
        }
    });

doubleIntegers.print();
> 2, 4, 6, 8


// Flat Map – Takes one element and produces zero, one, or more elements.
DataStream<Integer> doubleIntegers2 =

    integers.flatMap(new FlatMapFunction<Integer, Integer>() {
        @Override
        public void flatMap(Integer value, Collector<Integer> out) {
            out.collect(value * 2);
        }
    });

doubleIntegers2.print();
> 2, 4, 6, 8
```

Transformations: Filter

```java
// The DataStream
DataStream<Integer> integers = env.fromElements(1, 2, 3, 4);

DataStream<Integer> filtered =

    integers.filter(new FilterFunction<Integer>() {
        @Override
        public boolean filter(Integer value) {
            return value != 3;
        }
    });

filtered.print();
> 1, 2, 4
```
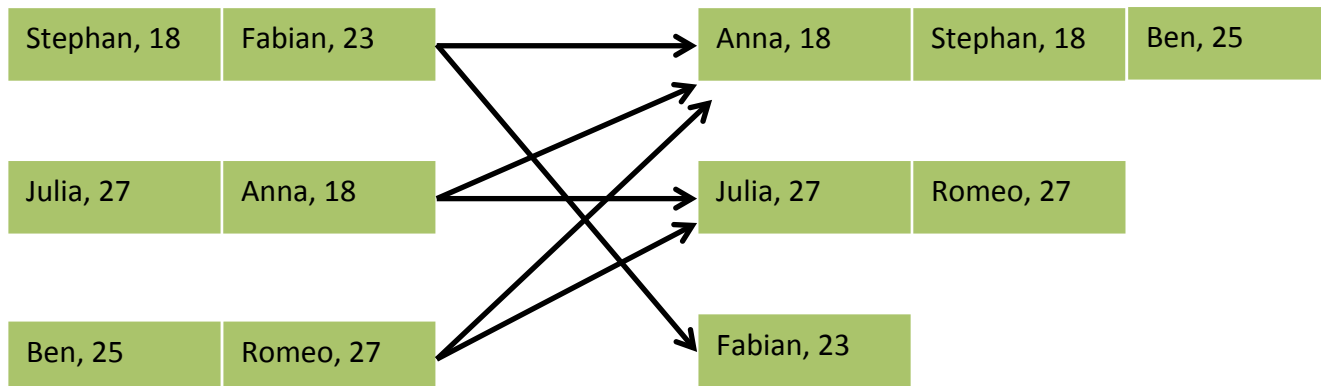
Transformations: KeyBy

- A DataStream can be organized by a key
  - Partitions the data (all elements with the same key are processed by the same operator)
  - Certain operators are key-aware
  - Operator state can be partitioned by key

```
// (name, age) of employees
DataStream<Tuple2<String, Integer>> passengers = …

// group by second field (age)
DataStream<Integer, Integer> grouped = passengers.keyBy(1)
```

Stephan, 18	Fabian, 23

Julia, 27	Anna, 18

Ben, 25	Romeo, 27

Anna, 18	Stephan, 18	Ben, 25

Julia, 27	Romeo, 27

Fabian, 23

Data Shipping Strategies

- Optionally, you can specify how data is shipped between two transformations

- Forward: `stream.forward()`
  - Only local communication

- Rebalance: `stream.rebalance()`
  - Round-robin partitioning

- Partition by hash: `stream.partitionByHash(...)`

- Custom partitioning: `stream.partitionCustom(...)`

- Broadcast: `stream.broadcast()`
  - Broadcast to all nodes

Rich Functions

RichFunctions

- Function interfaces have only one method
  - Single abstract method (SAM)
  - Support for Java8 Lambda functions

- There is a "Rich" variant for each function.
  - `RichFlatMapFunction`, …
  - Additional methods
    - `open(Configuration c)`
    - `close()`
    - `getRuntimeContext()`

RichFunctions & RuntimeContext

- RuntimeContext has useful methods:
    - `getIndexOfThisSubtask ()`
    - `getNumberOfParallelSubtasks()`
    - `getExecutionConfig()`

- Hands out partitioned state (later discussed)
    - `getKeyValueState()`

Data Sources

Data Sources: Collections

```
StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

// read from elements
DataStream<String> names = env.fromElements("Some", "Example",
"Strings");


// read from Java collection
List<String> list = new ArrayList<String>();
list.add("Some");
list.add("Example");
list.add("Strings");

DataStream<String> names = env.fromCollection(list);
```

Data Sources: Files & Sockets

```java
StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();


// read text socket from port
DataStream<String> socketLines = env
    .socketTextStream("localhost", 9999);


// read a text file ingesting new elements every 100 milliseconds
DataStream<String> localLines = env
    .readFileStream("/path/to/file", 100,
     WatchType.PROCESS_ONLY_APPENDED);
```

Custom SourceFunctions & Connectors

```
StreamExecutionEnvironment env =
    StreamExecutionEnvironment.getExecutionEnvironment();

// read data stream from custom source function
DataStream<<Tuple2<Long, String> stream = env
    .addSource(new MySourceFunction());
```

Flink has connectors for many stream serving systems

- Apache Kafka

- RabbitMQ

- ...

Data Sinks

Data Sinks

Write as text file using toString()
- `stream.writeAsText("/path/to/file")`

Write as CSV file
- `stream.writeAsCsv("/path/to/file")`

Print to the standard output
- `stream.print()`

Emit to socket
- `stream`
  `.writeToSocket(host,port,SerializationSchema)`

Custom Sinks & Connectors

Emit data with a custom sink function

- `stream.addSink(new MySinkFunction());`

Several connectors

- Apache Kafka
- Elasticsearch
- Rolling Files (HDFS, S3, …)

Execution

- Programs are lazily executed when `execute()` is called

```
DataStream<T> result;

// nothing happens
result.writeToSocket(...);

// nothing happens
result.writeAsText("/path/to/file", "\n", "|");

// Execution really starts here
env.execute();
```

Fault-Tolerance and Operator State

Fault Tolerance in Flink

- Flink takes a checkpoint of an application every *N* milliseconds and rolls back to the checkpointed state in case of a failure

- Enable exactly-once consistency (checkpoint every 5 seconds)
  ```
  env.enableCheckpointing(5000)
  ```

- Enable at-least-once consistency (for lower latency)
  ```
  env.enableCheckpointing(5000, CheckpointingMode.AT_LEAST_ONCE)
  ```

- Setting the interval to few seconds should be good for most application

- If checkpointing is not enabled, no recovery guarantees are provided

- See documentation for details:
  https://ci.apache.org/projects/flink/flink-docs-master/internals/stream_checkpointing.html

Stateful Functions

- All DataStream functions can be stateful
  - State is checkpointed and recovered in case of a failure (if checkpointing is enabled).

- You can define two types of state
  - *Local State*: Functions can register local variables to be checkpointed.
  - *Key-Partitioned State*: Functions on a keyed stream can access and update state scoped to the current key.

Defining Local State

```java
DataStream<String> aStream;
DataStream<Long> lengths = aStream.map(new MapWithCounter());

public static class MapWithCounter
    implements MapFunction<String, Long>, Checkpointed<Long> {

    private long totalLength = 0;

    @Override
    public Long map (String value) {
        totalLength += value.length();
        return totalLength;
    }

    @Override
    public Long snapshotState(long cpId, long cpTimestamp) throws Exception {
        return totalLength;
    }

    @Override
    public void restoreState(Long state) throws Exception {
        totalLength = state;
    }
}
```

Defining Key-Partitioned State

```java
DataStream<Tuple2<String, String>> aStream;
KeyedStream<Tuple2<String, String>, Tuple> keyedStream = aStream.keyBy(0);
DataStream<Long> lengths = keyedStream.map(new MapWithCounter());

public static class MapWithCounter
    extends RichMapFunction<Tuple2<String, String>,Long> {

    private OperatorState<Long> totalLengthByKey;

    @Override
    public void open (Configuration conf) {
        totalLengthByKey = getRuntimeContext()
                .getKeyValueState("totalLengthByKey", Long.class, 0L);
    }

    @Override
    public Long map (Tuple2<String,String> value) throws Exception {
        long newTotalLength = totalLengthByKey.value() + value.f1.length();
        totalLengthByKey.update(newTotalLength);
        return totalLengthByKey.value();
    }
}
```

38

Best Practices

Some advice

- Use `env.fromElements(..)` or `env.fromCollection(..)` to quickly get a DataStream to experiment with

- Use `print()` to print a DataStream