

DataStream API

Basics



Apache Flink® Training
dataArtisans

Flink v1.3 – 19.06.2017

DataStream API



- Stream Processing
- Java and Scala
- All examples here in Java for Flink 1.3
- Documentation available at
flink.apache.org

DataStream API by Example

Window WordCount: main Method



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        //sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Stream Execution Environment



```
public static void main(String[] args) throws Exception {  
    // set up the execution environment  
    final StreamExecutionEnvironment env =  
        StreamExecutionEnvironment.getExecutionEnvironment();  
  
    // configure event time  
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);  
  
    DataStream<Tuple2<String, Integer>> counts = env  
        // read stream of words from socket  
        .socketTextStream("localhost", 9999)  
        // split up the lines in tuples containing: (word,1)  
        .flatMap(new Splitter())  
        // key stream by the tuple field "0"  
        .keyBy(0)  
        // compute counts every 5 minutes  
        .timeWindow(Time.minutes(5))  
        // sum up tuple field "1"  
        .sum(1);  
  
    // print result in command line  
    counts.print();  
    // execute program  
    env.execute("Socket WordCount Example");  
}
```

Data Sources



```
public static void main(String[] args) throws Exception {  
    // set up the execution environment  
    final StreamExecutionEnvironment env =  
        StreamExecutionEnvironment.getExecutionEnvironment();  
  
    // configure event time  
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);  
  
    DataStream<Tuple2<String, Integer>> counts = env  
        // read stream of words from socket  
        .socketTextStream("localhost", 9999)  
        // split up the lines in tuples containing: (word,1)  
        .flatMap(new Splitter())  
        // key stream by the tuple field "0"  
        .keyBy(0)  
        // compute counts every 5 minutes  
        .timeWindow(Time.minutes(5))  
        // sum up tuple field "1"  
        .sum(1);  
  
    // print result in command line  
    counts.print();  
    // execute program  
    env.execute("Socket WordCount Example");  
}
```

Data types



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        //sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Transformations



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        // sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```


User functions



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        //sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

DataSinks



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        // sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Execute!



```
public static void main(String[] args) throws Exception {
    // set up the execution environment
    final StreamExecutionEnvironment env =
        StreamExecutionEnvironment.getExecutionEnvironment();

    // configure event time
    env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

    DataStream<Tuple2<String, Integer>> counts = env
        // read stream of words from socket
        .socketTextStream("localhost", 9999)
        // split up the lines in tuples containing: (word,1)
        .flatMap(new Splitter())
        // key stream by the tuple field "0"
        .keyBy(0)
        // compute counts every 5 minutes
        .timeWindow(Time.minutes(5))
        //sum up tuple field "1"
        .sum(1);

    // print result in command line
    counts.print();
    // execute program
    env.execute("Socket WordCount Example");
}
```

Window WordCount: FlatMap



```
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

    @Override
    public void flatMap(String value,
                        Collector<Tuple2<String, Integer>> out)
        throws Exception {
        // normalize and split the line
        String[] tokens = value.toLowerCase().split("\\W+");

        // emit the pairs
        for (String token : tokens) {
            if (token.length() > 0) {
                out.collect(
                    new Tuple2<String, Integer>(token, 1));
            }
        }
    }
}
```

WordCount: Interface



```
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

    @Override
    public void flatMap(String value,
                        Collector<Tuple2<String, Integer>> out)
        throws Exception {
        // normalize and split the line
        String[] tokens = value.toLowerCase().split("\\W+");

        // emit the pairs
        for (String token : tokens) {
            if (token.length() > 0) {
                out.collect(
                    new Tuple2<String, Integer>(token, 1));
            }
        }
    }
}
```

WordCount: Types



```
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

    @Override
    public void flatMap(String value,
                        Collector<Tuple2<String, Integer>> out)
        throws Exception {
        // normalize and split the line
        String[] tokens = value.toLowerCase().split("\\W+");

        // emit the pairs
        for (String token : tokens) {
            if (token.length() > 0) {
                out.collect(
                    new Tuple2<String, Integer>(token, 1));
            }
        }
    }
}
```

WordCount: Collector



```
public static class Splitter
    implements FlatMapFunction<String, Tuple2<String, Integer>> {

    @Override
    public void flatMap(String value,
                        Collector<Tuple2<String, Integer>> out)
        throws Exception {
        // normalize and split the line
        String[] tokens = value.toLowerCase().split("\\W+");

        // emit the pairs
        for (String token : tokens) {
            if (token.length() > 0) {
                out.collect(
                    new Tuple2<String, Integer>(token, 1));
            }
        }
    }
}
```

What kind of data can Flink handle?

DataStream API: Data Types

Data Types



- Flink aims to be able to process data of any type
- DataSet and DataStream APIs share the same type system
- Basic Types
 - String, Long, Integer, Boolean, ...
 - Arrays
- Composite Types
 - Tuples
 - POJOs
 - Scala Case Classes

Tuples



- Easiest and most efficient way to encapsulate data
- Scala: use default Scala tuples (1 to 22 fields)
- Java: Tuple1 up to Tuple25

```
Tuple2<String, String> person =  
    new Tuple2<>("Max", "Mustermann");
```

```
Tuple3<String, String, Integer> person =  
    new Tuple3<>("Max", "Mustermann", 42);
```

```
Tuple4<String, String, Integer, Boolean> person =  
    new Tuple4<>("Max", "Mustermann", 42, true);
```

```
// zero based index!  
String firstName = person.f0;  
String secondName = person.f1;  
Integer age = person.f2;  
Boolean fired = person.f3;
```

POJOs



- Any Java class that
 - Has an empty default constructor
 - Has publicly accessible fields
 - public field or default getter & setter

```
public class Person {  
    public int id;  
    public String name;  
    public Person() {};  
    public Person(int id, String name) {...};  
}
```

```
DataStream<Person> p =  
    env.fromElements(new Person(1, "Bob"));
```

Case Classes (Scala)



- Scala case classes are natively supported

```
case class Person(id: Int, name: String)  
  
d: DataStream[Person] =  
    env.fromElements(Person(1, "Bob"))
```

DataStream API: Operators

Transformations: map & flatMap



```
DataStream<Integer> integers = env.fromElements(1, 2, 3, 4);
```

```
// Regular Map - Takes one element and produces one element
```

```
DataStream<Integer> doubleIntegers =  
    integers.map(new MapFunction<Integer, Integer>() {  
        @Override  
        public Integer map(Integer value) {  
            return value * 2;  
        }  
    });
```

```
doubleIntegers.print();
```

```
> 2, 4, 6, 8
```

```
// Flat Map - Takes one element and produces zero, one, or more elements
```

```
DataStream<Integer> doubleIntegers2 =  
    integers.flatMap(new FlatMapFunction<Integer, Integer>() {  
        @Override  
        public void flatMap(Integer value, Collector<Integer> out) {  
            out.collect(value * 2);  
        }  
    });
```

```
doubleIntegers2.print();
```

```
> 2, 4, 6, 8
```

Transformations: Filter



```
// The DataStream
DataStream<Integer> integers = env.fromElements(1, 2, 3, 4);

DataStream<Integer> filtered =
    integers.filter(new FilterFunction<Integer>() {
        @Override
        public boolean filter(Integer value) {
            return value != 3;
        }
    });

filtered.print();
> 1, 2, 4
```


Transformations: KeyBy



- A `DataStream` can be organized by a key
 - Partitions the data, i.e., all elements with the same key are processed by the same operator
 - Certain operators are key-aware
 - Operator state can be partitioned by key

```
// (name, age) of passengers  
DataStream<Tuple2<String, Integer>> passengers = ...
```

```
// key by second field (age)  
DataStream<Tuple2<String, Integer>> grouped = passengers.keyBy(1)
```



Reduce (conceptually)



```
public Integer reduce(Integer a, Integer b) {  
    return a + b;  
}
```

[1, 2, 3, 4] → **reduce**() means: (((1 + 2) + 3) + 4) = 10

Reduce on a Stream



- Can only be used with keyed or windowed streams
- Example with reduce on a KeyedStream

```
// Produce running sums of the even and odd integers.
```

```
List  data = new ArrayList
```

```
data.add(new Tuple2<>("odd", 1));  
data.add(new Tuple2<>("even", 2));  
data.add(new Tuple2<>("odd", 3));  
data.add(new Tuple2<>("even", 4));
```

```
DataStream  tuples = env.fromCollection(data);
```

```
KeyedStream  odd_and_evens = tuples.keyBy(0);
```

Reduce on a KeyedStream



```
DataStream<Tuple2<String, Integer>> sums =  
    odd_and_evens.reduce(new ReduceFunction<Tuple2<String, Integer>>() {  
        @Override  
        public Tuple2<String, Integer> reduce(  
            Tuple2<String, Integer> t1,  
            Tuple2<String, Integer> t2) throws Exception {  
                return new Tuple2<>(t1.f0, t1.f1 + t2.f1);  
            }  
        });
```

```
sums.print();  
env.execute();
```

```
3> (odd,1)  
3> (odd,4)  
4> (even,2)  
4> (even,6)
```

Specifying Keys

Keyed Streams



- `keyBy()` partitions a `DataStream`
 - a key is extracted from each element
- Basis of operating in parallel
- Data types used as keys must have valid implementations of `hashCode()` and `equals()`, which rules out arrays (for example)
- Composite types can be used as keys
 - all the fields must be key types
 - nested fields can also be used as keys

Keys for Tuples



- Define keys by field position

```
DataStream<Tuple3<Integer, String, Double>> d = ...
```

```
// key stream by String field  
d.keyBy(1);
```

- Or field names

```
// key stream by Double field  
d.keyBy("f2");
```

Keys for POJOs



- Define keys by field name

```
DataStream<Person> d = ...
```

```
// key stream by “name” field  
d.keyBy("name");
```

Keys for Case Classes (Scala)



- Define keys by field name

```
case class Person(id: Int, name: String)  
d: DataStream[Person] = ...
```

```
// key stream by field "name"  
d.keyBy("name")
```


Working With Multiple Streams

Connected Streams



- Connect two streams to correlate them with each other
- Apply functions on connected streams to share state
- Typical use case is to use one stream for control and another for data

```
DataStream<String> control = ...  
DataStream<String> data = ...
```

```
DataStream<String> result = control  
    .connect(data)  
    .flatMap(new MyCoFlatMap());
```

FlatMap on Connected Streams



```
private static final class MyCoFlatMap
    implements CoFlatMapFunction<String, String, String> {
    HashSet blacklist = new HashSet();

    @Override
    public void flatMap1(String control_value, Collector<String> out) {
        blacklist.add(control_value);
        out.collect("listed " + control_value);
    }

    @Override
    public void flatMap2(String data_value, Collector<String> out) {
        if (blacklist.contains(data_value)) {
            out.collect("skipped " + data_value);
        } else {
            out.collect("passed " + data_value);
        }
    }
}
```

FlatMap on Connected Streams



```
StreamExecutionEnvironment env =  
    StreamExecutionEnvironment.getExecutionEnvironment();  
  
DataStream<String> control =  
    env.fromElements("DROP", "IGNORE");  
  
DataStream<String> data =  
    env.fromElements("data", "DROP", "artisans", "IGNORE");  
  
DataStream<String> result = control  
    .broadcast()  
    .connect(data)  
    .flatMap(new MyCoFlatMap());  
  
result.print();  
env.execute();
```

FlatMap on Connected Streams



```
control = env.fromElements("DROP", "IGNORE");  
data = env.fromElements("data", "DROP", "artisans", "IGNORE");
```

```
...
```

```
env.execute();
```

```
> listed DROP  
> listed IGNORE  
> passed data  
> skipped DROP  
> passed artisans  
> skipped IGNORE
```

Broadcast



```
DataStream<String> result = control  
    .broadcast()  
    .connect(data)  
    .flatMap(new MyCoFlatMap());
```

- Events are replicated to all downstream operators
- This is not a magical, managed, replicated state solution
- And you have to consider the race condition implications

Can also use Map on Connected Streams



```
DataStream<String> strings = ...
```

```
DataStream<Integer> ints = ...
```

```
ints.connect(strings)
```

```
    .map(new CoMapFunction<Integer, String, Boolean>() {
```

```
        @Override
```

```
        public Boolean map1 (Integer value) {
```

```
            return true;
```

```
        }
```

```
        @Override
```

```
        public Boolean map2 (String value) {
```

```
            return false;
```

```
        }
```

```
    });
```

Rich Functions

Rich Functions



- Function interfaces have only one method
 - Single abstract method (SAM)
 - Support for Java8 lambda functions

- There is a “Rich” variant of each function type
 - `RichFlatMapFunction`, ...
 - Additional methods
 - `open(Configuration c)`
 - `close()`
 - `getRuntimeContext()`

Rich Functions & RuntimeContext



- RuntimeContext has useful methods
 - `getIndexOfThisSubtask()`
 - `getNumberOfParallelSubtasks()`
 - `getExecutionConfig()`

- RuntimeContext also provides access to partitioned state (discussed later)
 - `getState()`

Wrap-up

Some tips



- Use `env.fromElements(...)` or `env.fromCollection(...)` to quickly create a `DataStream` to experiment with
- Use `print()` to print a `DataStream`
- Lazy execution can make debugging tricky, but you can use breakpoints in your IDE