

DataStream API

State & Failure Recovery



Apache Flink® Training

dataArtisans

Flink v1.3 – 20.06.2017

Working with (Rescalable) State

Stateful Functions



- All DataStream functions can be stateful
 - Flink manages state so that it can be redistributed/rescaled
 - State is checkpointed and restored in case of a failure (if checkpointing is enabled)


- Flink manages two types of state:
 - Operator (non-keyed) state
 - Keyed state

- Flink supports rescaling the state it manages

Operator vs Keyed State




Operator (non-keyed)




- State bound only to operator
- E.g. source state

Keyed




- State bound to an operator + key
- E.g. Keyed UDF and window state
- "SELECT count(*) FROM t GROUP BY t.key"

Repartitioning Operator State



Operator state: a list of state elements which can be freely repartitioned

Scaling out





Repartitioning Keyed State



- Split key space into key groups
- Every key falls into exactly one key group
- Assign key groups to tasks
- Maximum parallelism defined by #key groups

Key space





Types of Keyed State



- `ValueState<T>`
- `ListState<T>`
- `ReducingState<T>`
- `MapState<UK, UV>` (*new in 1.3*)
- ~~`FoldingState<T>`~~ (*deprecated in 1.3*)
 - `AggregatingState<IN, OUT>`

Using Key-Partitioned State



```
DataStream<Tuple2<String, String>> strings = ...
DataStream<Long> lengths = strings
    .keyBy(0)
    .map(new MapWithCounter());
```

```
public static class MapWithCounter extends RichMapFunction<Tuple2<String, String>, Long> {
    // state object
    private ValueState<Long> totalLengthByKey;

    @Override
    public void open (Configuration conf) {
        // obtain state object
        ValueStateDescriptor<Long> descriptor = new ValueStateDescriptor<>(
            "totalLengthByKey", Long.class);
        totalLengthByKey = getRuntimeContext().getState(descriptor);
    }

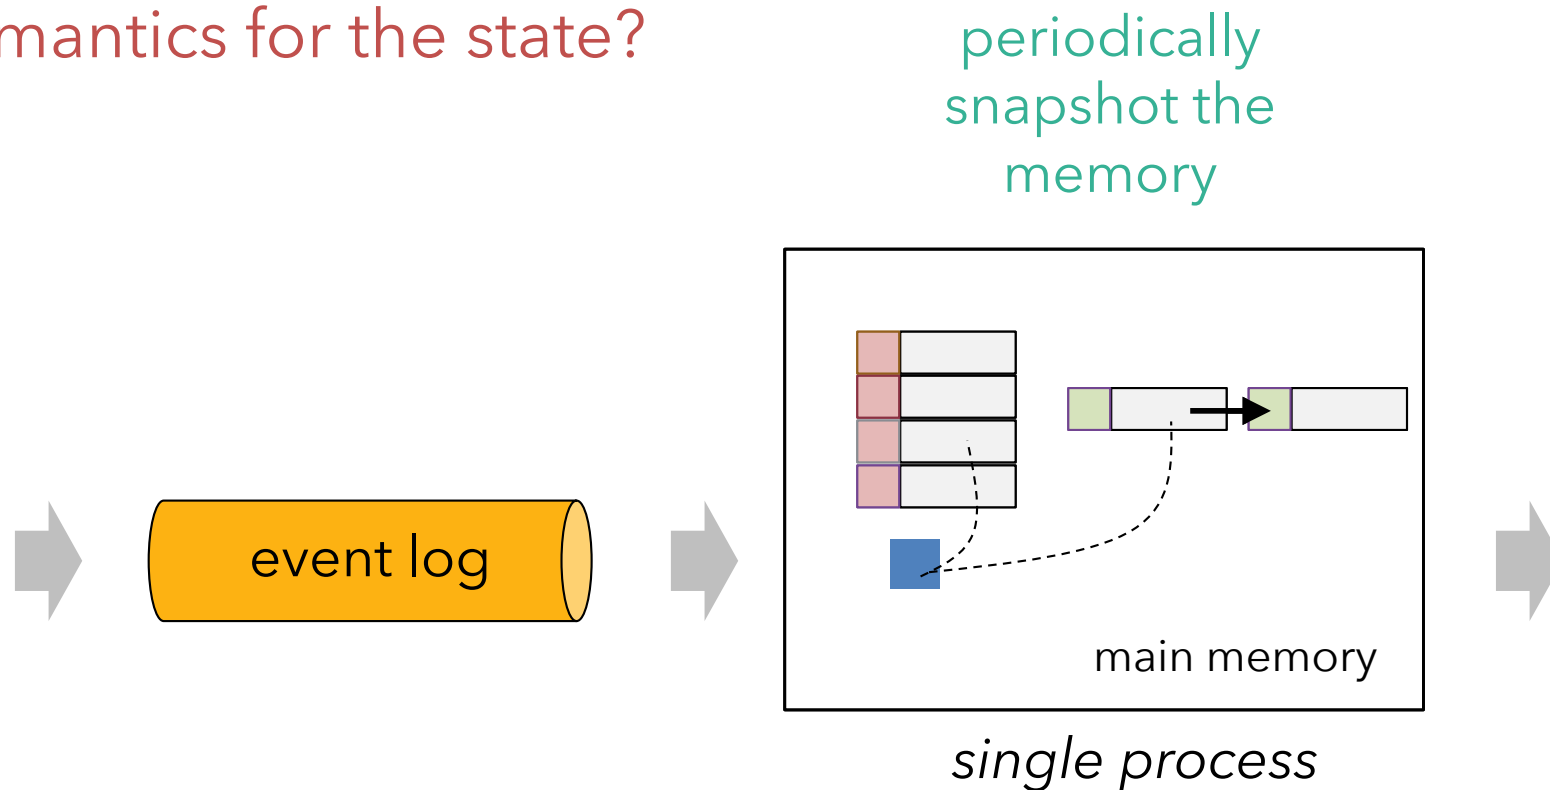
    @Override
    public Long map (Tuple2<String, String> value) throws Exception {
        long length = totalLengthByKey.value();    // fetch state for current key
        if (length == null) length = 0;
        long newTotalLength = length + value.f1.length();
        totalLengthByKey.update(newTotalLength);    // update state of current key
        return totalLengthByKey.value();
    }
}
```

Fault Tolerance via Snapshotting

Fault tolerance: simple case




How to ensure exactly-once semantics for the state?



Fault tolerance: simple case



Recovery: restore
snapshot and replay
events since snapshot



State fault tolerance



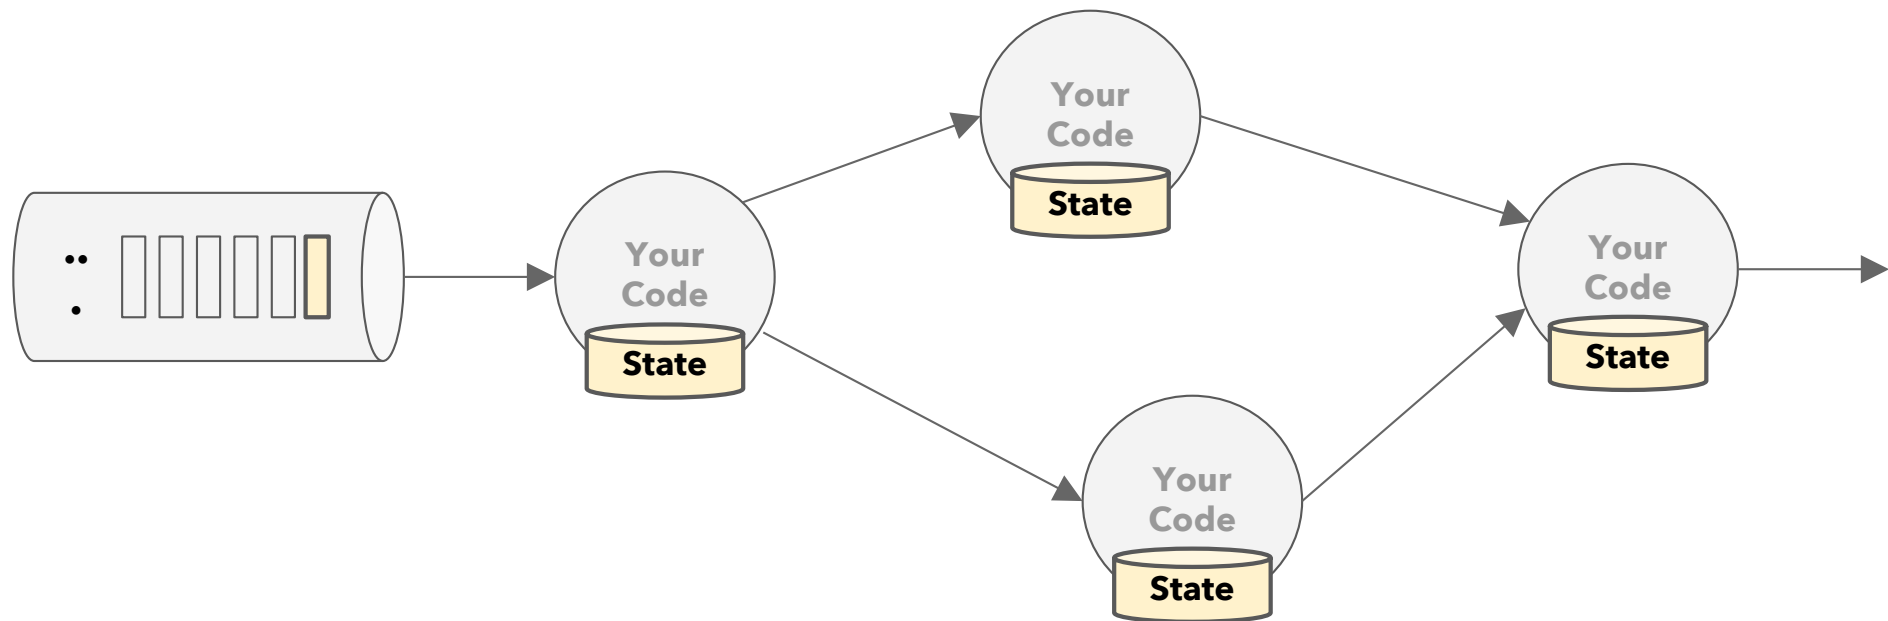
Fault tolerance concerns for a stateful stream processor:

- How to ensure exactly-once semantics for the state?
- How to create consistent snapshots of **distributed embedded state**?
- More importantly, how to do it **efficiently without interrupting computation**?

State fault tolerance (II)



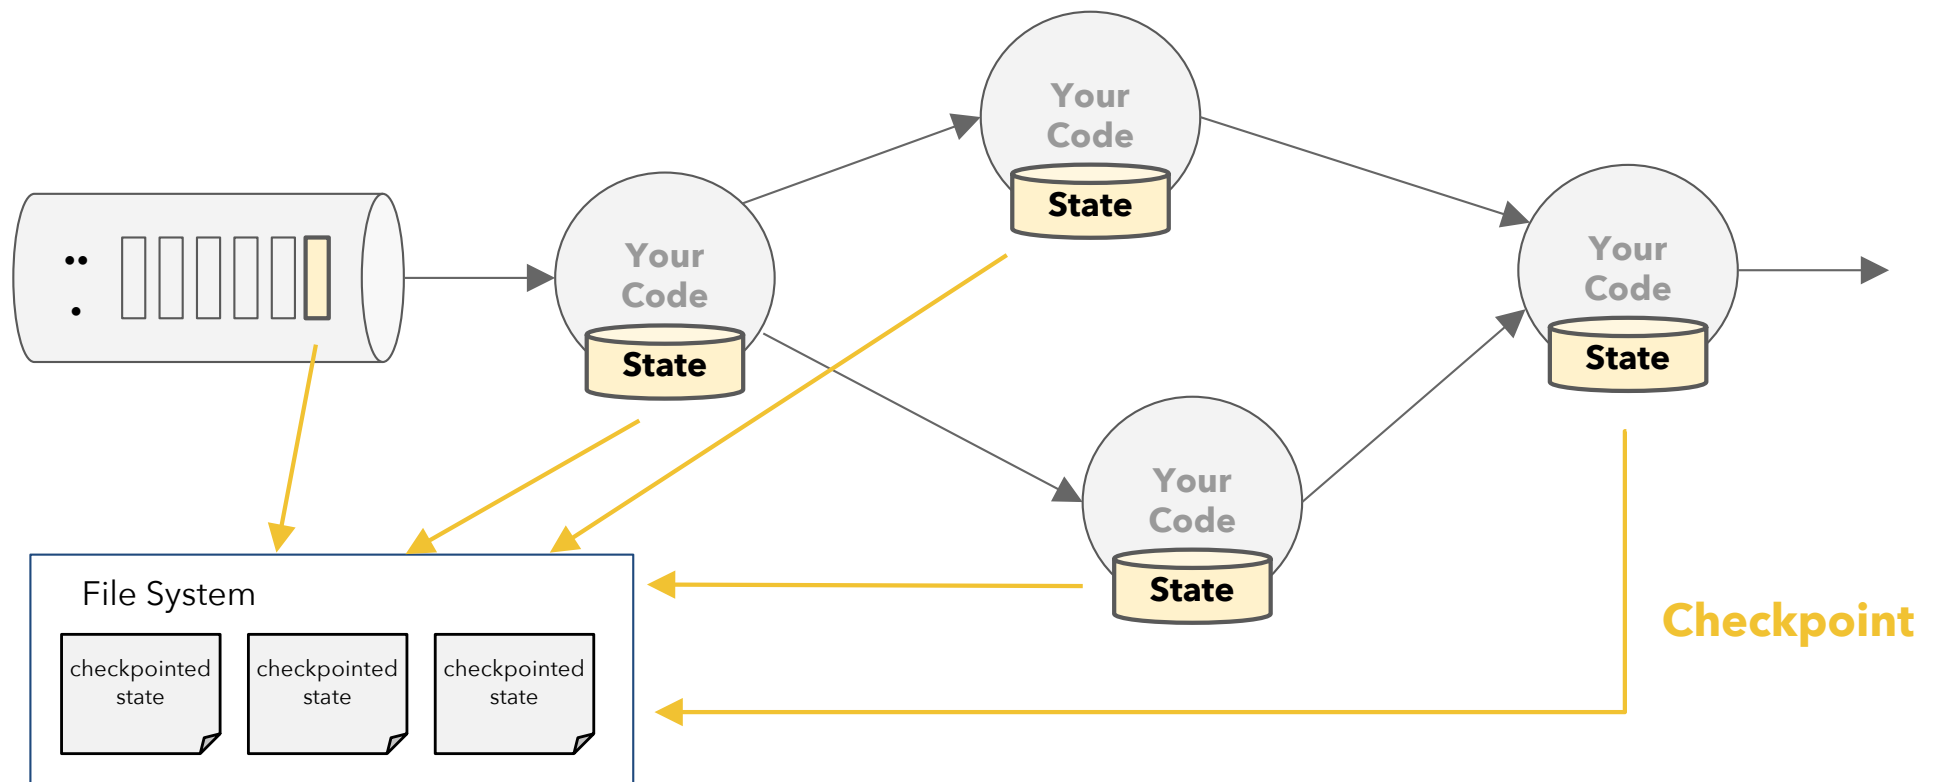
Consistent snapshotting:



State fault tolerance (II)



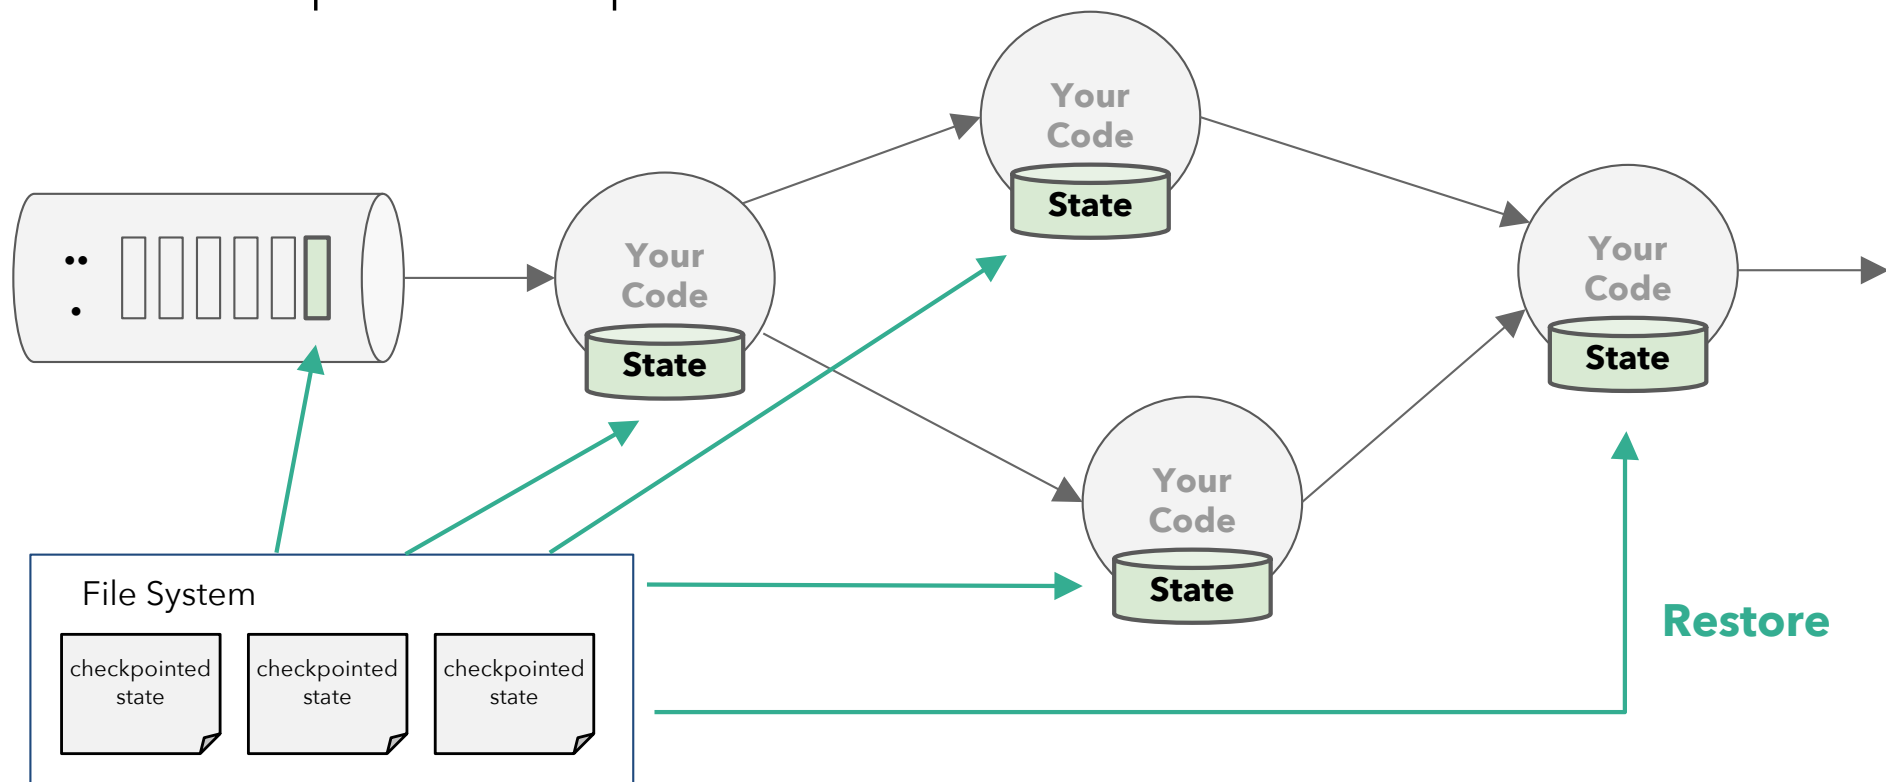
Consistent snapshotting:



State fault tolerance (III)



- Recover all embedded state
- Reset position in input stream



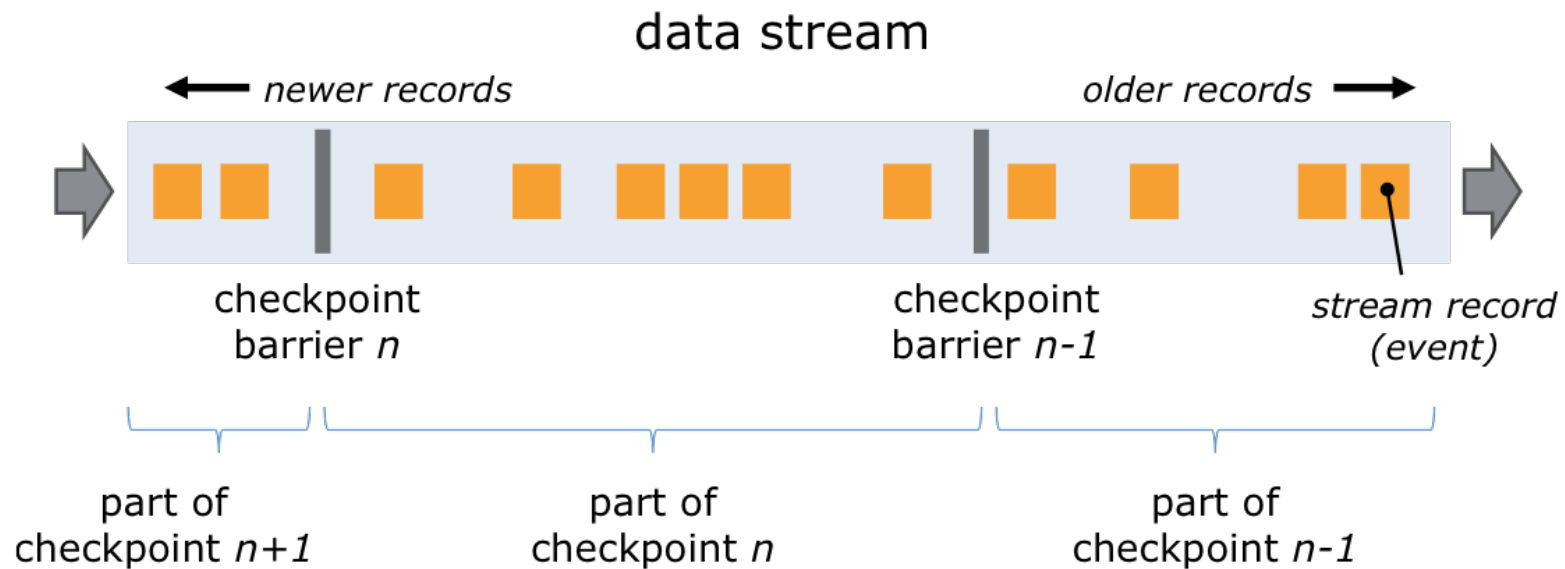
Checkpoints

Checkpointing in Flink




- Asynchronous Barrier Snapshotting
 - checkpoint barriers are inserted into the stream and flow through the graph along with the data
 - this avoids a "global pause" during checkpointing
- Checkpoint barriers cause ...
 - replayable sources to checkpoint their offsets
 - operators to checkpoint their state
 - sinks to commit open transactions
- The program is rolled back to the latest completed checkpoint in case of a failure.


Checkpoint Barriers




Asynchronous Barrier Snapshotting




begin aligning



aligning



checkpoint




continue


Distributed snapshots




Events flow without replication or synchronous writes



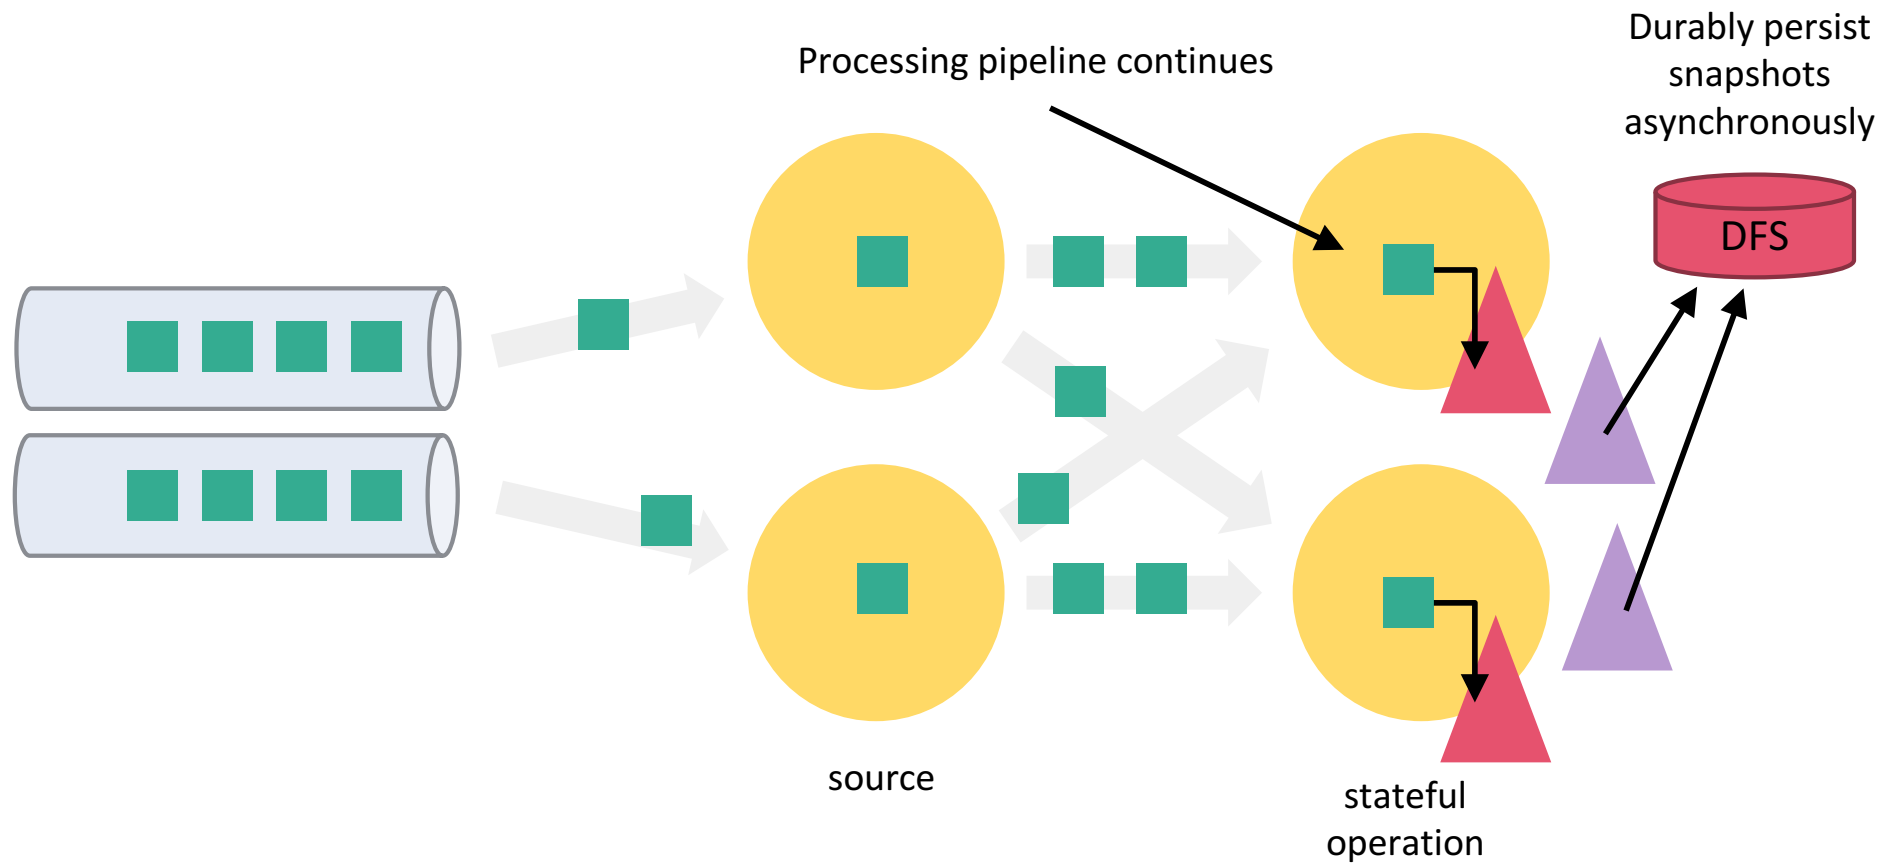
Distributed snapshots



Distributed snapshots



Distributed snapshots



Enabling Checkpointing



- Checkpointing is disabled by default.
- Enable checkpointing with exactly once consistency:

```
// checkpoint every 5 seconds  
env.enableCheckpointing(5000)
```

- Configure at least once consistency (for lower latency):

```
env.getCheckpointConfig()  
    .setCheckpointingMode(CheckpointingMode.AT_LEAST_ONCE);
```

- Most applications perform well with a few seconds checkpointing interval.

Restart Strategies



- How often and fast does a job try to restart?
- Available strategies
 - No restart (default)
 - Fixed delay
 - Failure rate

```
// Fixed Delay restart strategy
env.setRestartStrategy(
  RestartStrategies.fixedDelayRestart(
    3, // no of restart attempts
    Time.of(10, TimeUnit.SECONDS) // restart interval
  ));
```

State Backends

State in Flink



- There are several sources of state in Flink
 - Windows
 - User functions
 - Sources and Sinks
 - Timers
- State is persisted during checkpoints, if checkpointing is enabled
- Internal representation and storage location depend on the configured State Backend

Choosing a State Backend



Name	Working state	State backup	Snapshotting
RocksDBStateBackend	Local disk (tmp directory)	Distributed file system	Asynchronously
<ul style="list-style-type: none">• Good for state larger than available memory• 10x slower than memory-based backends			
FsStateBackend	JVM Heap	Distributed file system	Synchronous / Async option in Flink 1.3
<ul style="list-style-type: none">• Fast, requires large heap			
MemoryStateBackend	JVM Heap	JobManager JVM Heap	Synchronous / Async option in Flink 1.3
<ul style="list-style-type: none">• Good for testing and experimentation with small state (locally)			

State Backend Configuration



- Configuration of default state backend in

`./conf/flink-conf.yaml`

- State backend configuration in job

```
env.setStateBackend(  
    new FsStateBackend(  
        "hdfs://namenode:40010/flink/checkpoints"  
    ));
```

Savepoints

State management



Two important management concerns for a long-running job:

- Can I change / fix bugs in my streaming pipeline? How do I handle job downtime?
- Can I rescale (change parallelism of) my computation?

State management: Savepoints



- A persistent snapshot of all state
- When starting an application, state can be initialized from a savepoint
- In-between savepoint and restore we can update Flink version or user code

Savepoints

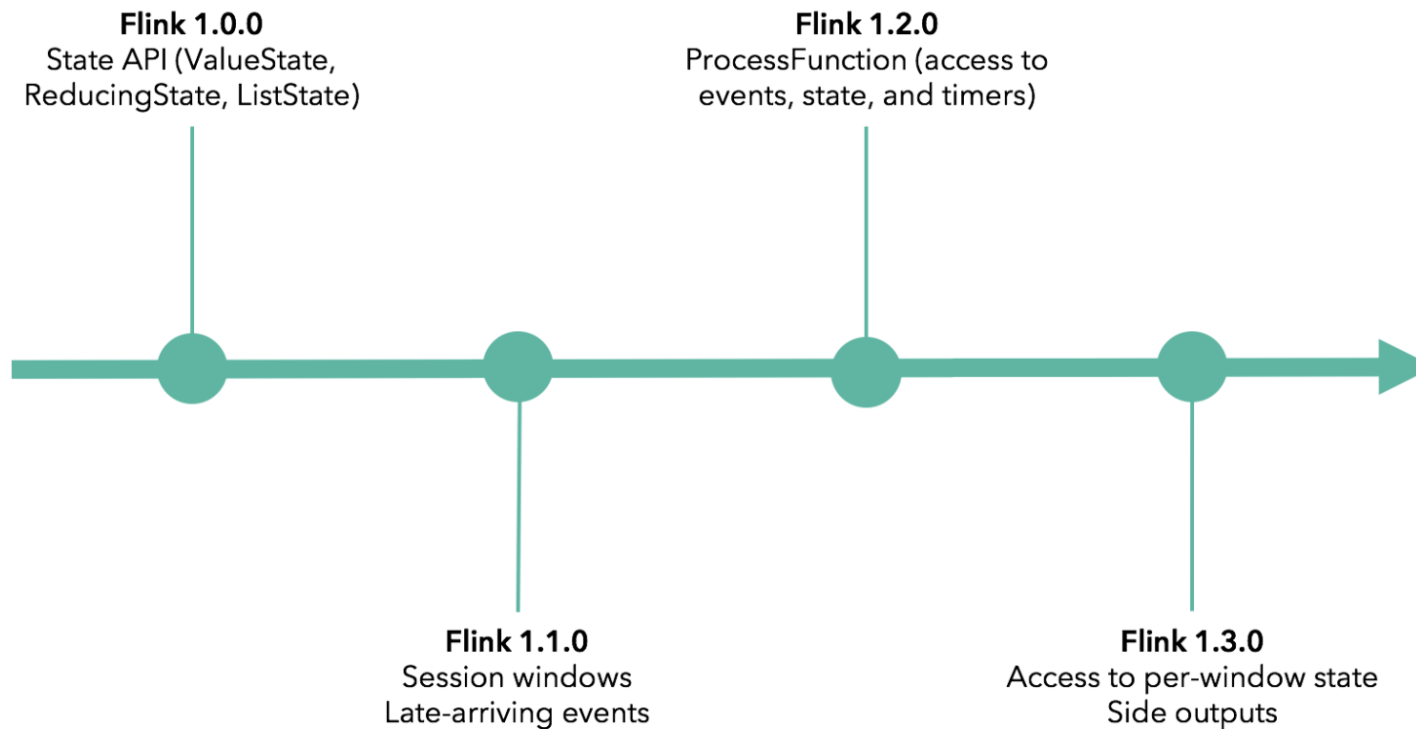


- A *Checkpoint* is a globally consistent point-in-time snapshot of a streaming application (*point in stream, state*)
- *Savepoints* are user-triggered, retained checkpoints
- Rescaling (currently) requires restarting from a savepoint
- Currently, Flink can only restore to the same state backend that created the savepoint

Evolution of Flink

(w.r.t. Stateful Stream Processing)

Evolution of Programming APIs



Evolution of Large State Handling

