

DataStream API

Windows & Time



Apache Flink® Training

dataArtisans

Flink v1.1.3 – 8.11.2016

Windows and Aggregates

Windows




- Aggregations on DataStreams are different from aggregations on DataSets
 - You cannot count all records of an infinite stream
- DataStream aggregations make sense on windowed streams
 - A finite subset of stream elements

Tumbling Windows




Aligned, fixed length, non-overlapping windows.



Sliding Windows




Aligned, fixed length, overlapping windows.



Session Windows



Non-aligned, variable length windows.



Specifying Windowing



```
// (name, age) of passengers
DataStream<Tuple2<String, Integer>> passengers = ...

passengers
    // group by second field (age)
    .keyBy(1)
    // window definition: tumbling window of 1 minute
    .timeWindow(Time.minutes(1))
```

Predefined Keyed Windows



- Tumbling time window
`.timeWindow(Time.minutes(1))`
- Sliding time window
`.timeWindow(Time.minutes(1), Time.seconds(10))`
- Tumbling count window
`.countWindow(100)`
- Sliding count window
`.countWindow(100, 10)`
- Session window
`.window(SessionWindows.withGap(Time.minutes(30)))`

Predefined Non-keyed Windows



- Windows on non-keyed streams are not processed in parallel!
- TimeWindow (tumbling, 10 seconds)
`.timeWindowAll(Time.seconds(10))`
- CountWindow (sliding, 20/10)
`.countWindowAll(20, 10)`

Aggregations on Windowed Streams



```
// (name, age) of passengers
DataStream<Tuple2<String, Integer>> passengers = ...

passengers
    // group by second field (age)
    .keyBy(1)
    // windows that are 1 minute long
    .timeWindow(Time.minutes(1))
    // apply a custom window function on window data
    .apply(new CountByAge());
```

Aggregation with a WindowFunction



```
public static class CountByAge implements WindowFunction<
    Tuple2<String, Integer>,           // input type
    Tuple3<Integer, Long, Integer>,    // output type
    Tuple,                             // key type
    TimeWindow> {                     // window type

    @Override
    public void apply(
        Tuple key,
        TimeWindow window,
        Iterable<Tuple2<String, Integer>> persons,
        Collector<Tuple3<Integer, Long, Integer>> out) {

        int age = ((Tuple1<Integer>)key).f0;
        int cnt = 0;

        for (Tuple2<String, Integer> p : persons) {
            cnt++;
        }
        out.collect(new Tuple3<>(age, window.getEnd(), cnt));
    }
}
```

Window State during Aggregation




state



Window State during Aggregation



state



Window State during Aggregation




state



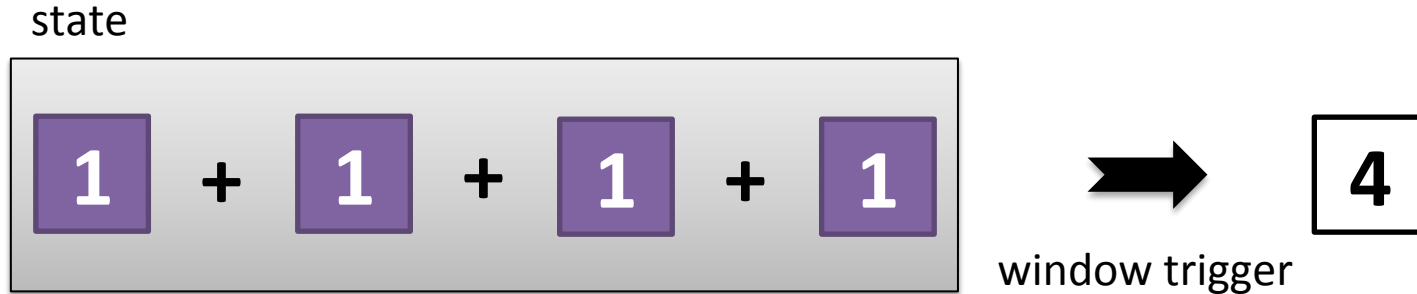
Window State during Aggregation



state



Window State during Aggregation



Operations on Windowed Streams



- `reduce(reduceFunction)`
 - Apply a functional reduce function to the window
- `fold(initialVal, foldFunction)`
 - Apply a functional fold function with a specified initial value to the window
- Aggregation functions
 - `sum()`, `min()`, `max()`, and others

Incremental Aggregation



state



Incremental Aggregation



state

2

Incremental Aggregation



state

3

Incremental Aggregation



state

4

Incremental Aggregation



state



window trigger

Incremental Window Aggregation



```
DataStream<Tuple2<String, Integer>> passengers = ...
```

```
passengers
    .keyBy(1)
    .timeWindow(Time.minutes(1), Time.seconds(10))
    .apply(new Tuple3<Integer, Long, Integer>(0, 0L, 0), new MyFoldFunction(), new MyWindowFunction())
```

```
private static class MyFoldFunction
    implements FoldFunction<Tuple2<String, Integer>, Tuple3<Integer, Long, Integer>> {

    public Tuple3<Integer, Long, Integer> fold(Tuple3<Integer, Long, Integer> acc, Tuple2<String, Integer> p) {
        Integer count = acc.getField(2);
        acc.setField(2, count + 1);
        return acc;
    }
}
```

```
private static class MyWindowFunction
    implements WindowFunction<Tuple3<Integer, Long, Integer>, Tuple3<Integer, Long, Integer>, Integer, TimeWindow> {

    public void apply(Integer age_key,
        TimeWindow window,
        Iterable<Tuple3<Integer, Long, Integer>> counts,
        Collector<Tuple3<Integer, Long, Integer>> out) {
        Integer count = counts.iterator().next().getField(2);
        out.collect(new Tuple3<Integer, Long, Integer>(age_key, window.getEnd(), count));
    }
}
```

Incremental Window Aggregation



```
DataStream<Tuple2<String, Integer>> passengers = ...
```

```
passengers
    .keyBy(1)
    .timeWindow(Time.minutes(1), Time.seconds(10))
    .apply(new Tuple3<Integer, Long, Integer>(0, 0L, 0), new MyFoldFunction(), new MyWindowFunction())
```

```
private static class MyFoldFunction
```

```
    implements FoldFunction<Tuple2<String, Integer>, Tuple3<Integer, Long, Integer>> {
```

```
    public Tuple3<Integer, Long, Integer> fold(Tuple3<Integer, Long, Integer> acc, Tuple2<String, Integer> p) {
        Integer count = acc.getField(2);
        acc.setField(2, count + 1);
        return acc;
    }
}
```

```
private static class MyWindowFunction
```

```
    implements WindowFunction<Tuple3<Integer, Long, Integer>, Tuple3<Integer, Long, Integer>, Integer, TimeWindow> {
```

```
    public void apply(Integer age_key,
        TimeWindow window,
        Iterable<Tuple3<Integer, Long, Integer>> counts,
        Collector<Tuple3<Integer, Long, Integer>> out) {
        Integer count = counts.iterator().next().getField(2);
        out.collect(new Tuple3<Integer, Long, Integer>(age_key, window.getEnd(), count));
    }
}
```

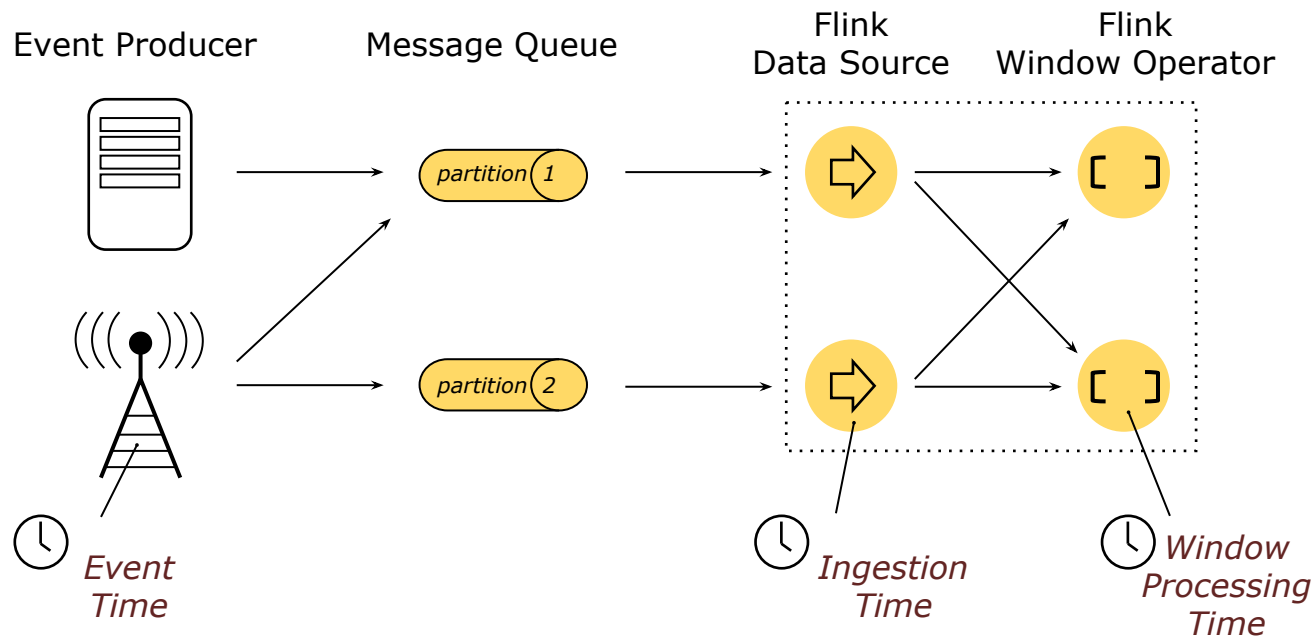

Custom window logic



- The DataStream API allows you to define very custom window logic
- GlobalWindows
 - a flexible, low-level window assignment scheme that can be used to implement custom windowing behaviors
 - only useful if you explicitly specify triggering, otherwise nothing will happen
- Trigger
 - defines when to evaluate a window
 - whether to purge the window or not
- *Careful!* This part of the API requires a good understanding of the windowing mechanism!

Handling Time Explicitly

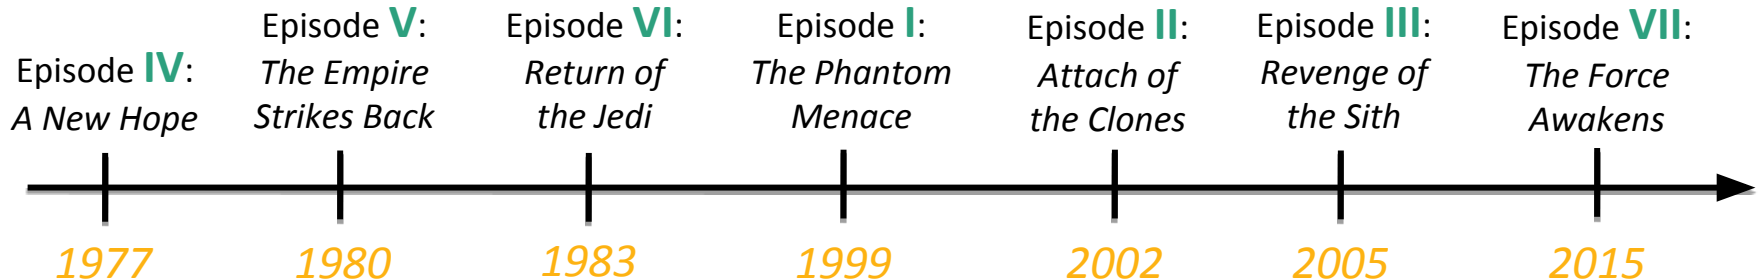
Different Notions of Time



Event Time vs Processing Time



This is called **event time**



This is called **processing time**

Setting the StreamTimeCharacteristic



```
final StreamExecutionEnvironment env =  
    StreamExecutionEnvironment.getExecutionEnvironment();  
  
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);  
  
// alternatively:  
// env.setStreamTimeCharacteristic(TimeCharacteristic.IngestionTime);  
// env.setStreamTimeCharacteristic(TimeCharacteristic.ProcessingTime);
```

Choosing Event Time has Consequences

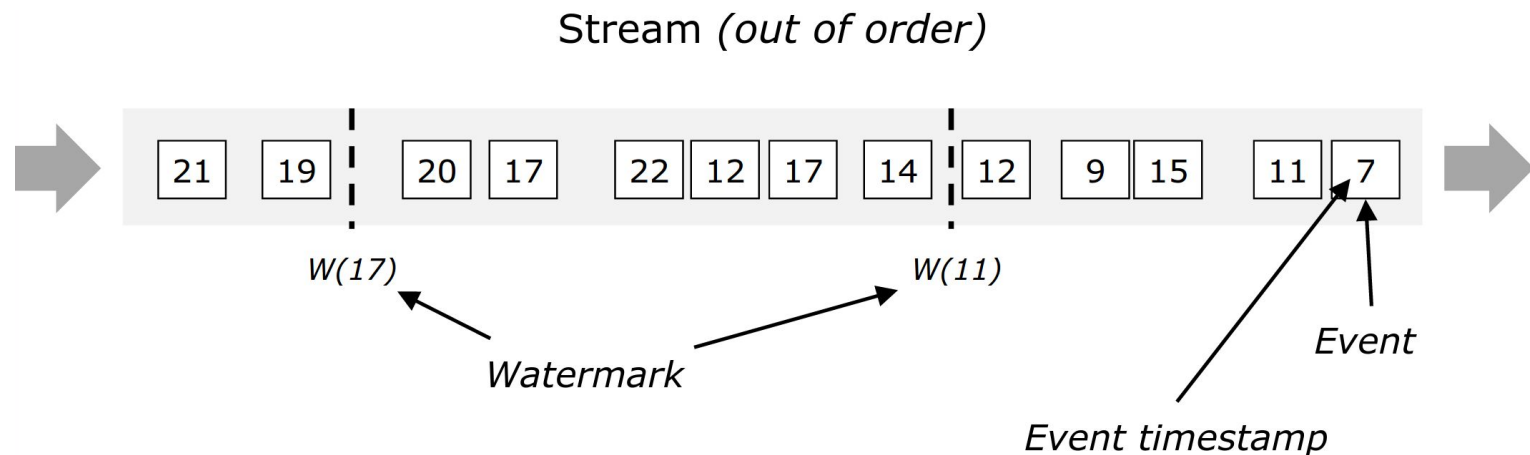


- When working with event time, Flink needs to know
 - how to extract the timestamp from a stream element
 - when enough event time has elapsed that a time window should be triggered

Watermarks



- Watermarks flow with the data stream and carry a timestamp; they are crucial for handling out-of-order events
- A $Watermark(t)$ is a declaration that all events with a $timestamp < t$ have occurred



Timestamp Assigners / Watermark Generators



```
DataStream<MyEvent> stream = ...
```

```
DataStream<MyEvent> withTimestampsAndWatermarks = stream  
    .assignTimestampsAndWatermarks(new MyTSExtractor());
```

```
withTimestampsAndWatermarks  
    .keyBy(...)  
    .timeWindow(...)  
    .addSink(...);
```




- There are different types of timestamp extractors
- BoundedOutOfOrdernessTimestampExtractor
 - Periodically emits watermarks that lag a fixed amount of time behind the max timestamp seen so far
 - To use, subclass and implement
 - `public abstract long extractTimestamp(T element)`
 - Constructor
 - `public BoundedOutOfOrdernessTimestampExtractor(
Time maxOutOfOrderness)`

References



- Documentation

- https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/event_time.html
- https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/event_timestamps_watermarks.html
- <https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/windows.html>

- Blog posts

- <http://flink.apache.org/news/2015/12/04/Introducing-windows.html>
- <http://data-artisans.com/how-apache-flink-enables-new-streaming-applications-part-1/>
- <https://www.mapr.com/blog/essential-guide-streaming-first-processing-apache-flink>
- <http://data-artisans.com/session-windowing-in-flink/>