

DataStream API

Windows & Time



Apache Flink® Training

dataArtisans

Flink v1.3 – 19.06.2017

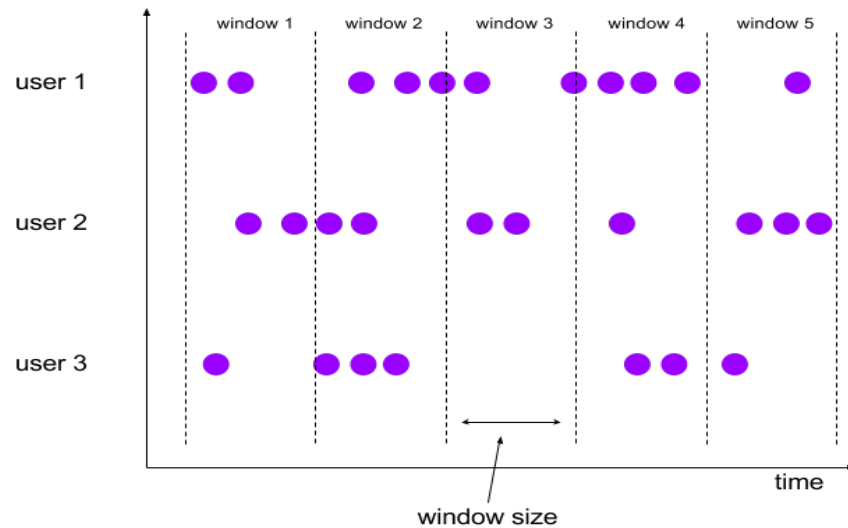
Windows and Aggregates

Windows

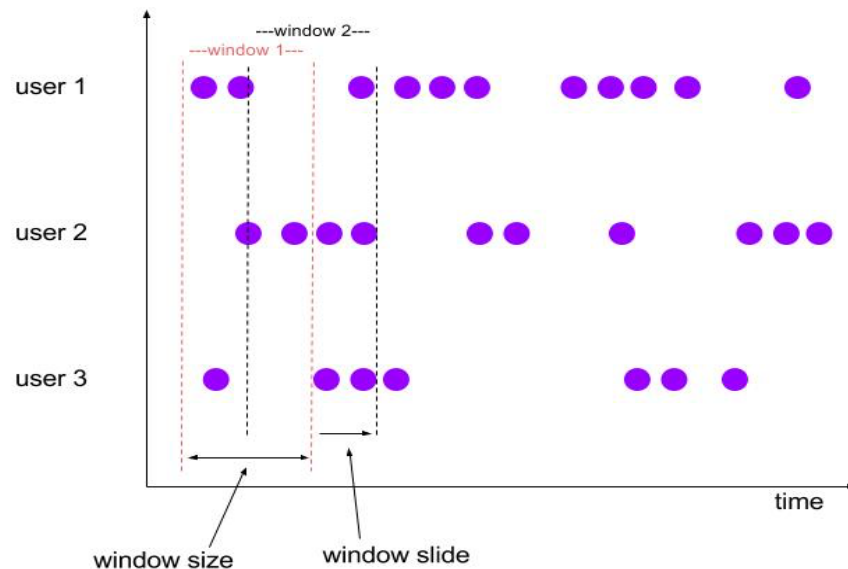


- Aggregations on streams are different from aggregations on batched data
 - You cannot count all records of an unbounded stream
- Aggregations do make sense on windowed streams, e.g.,
 - Number of transactions per minute

Tumbling and Sliding Windows



Tumbling:
aligned, fixed length,
non-overlapping windows




Sliding:
aligned, fixed length,
overlapping windows

Session Windows



Non-aligned, variable length windows.



Specifying Windowing



stream

<code>.keyBy(...)</code>	/ keyed vs non-keyed windows
<code>.window(...)</code>	/ “Assigner”
<code>.trigger(...)</code>	/ each Assigner has a default Trigger
<code>.evictor(...)</code>	/ default: no Evictor
<code>.allowedLateness()</code>	/ default: zero
<code>.process apply reduce()</code>	/ window function

Predefined Keyed Windows



- **Tumbling time window**
`.timeWindow(Time.minutes(1))`
- **Sliding time window**
`.timeWindow(Time.minutes(1), Time.seconds(10))`
- **Tumbling count window**
`.countWindow(100)`
- **Sliding count window**
`.countWindow(100, 10)`
- **Session window**
`.window(SessionWindows.withGap(Time.minutes(30)))`

Non-keyed Windows



- Windows on non-keyed streams are not processed in parallel!
 - `stream.windowAll(...)`
 - `stream.timeWindowAll(Time.seconds(10))...`
 - `stream.countWindowAll(20, 10)...`

Aggregations on Windowed Streams



```
DataStream<SensorReading> input = ...
```

```
input
    .keyBy("key")
    .timeWindow(Time.minutes(1))
    .apply(new MyWastefulMax());
```

```
public static class MyWastefulMax implements WindowFunction<
    SensorReading,                      // input type
    Tuple3<String, Long, Integer>,      // output type
    Tuple,                             // key type
    TimeWindow> {                      // window type

    @Override
    public void apply(
        Tuple key,
        TimeWindow window,
        Iterable<SensorReading> events,
        Collector<Tuple3<String, Long, Integer>> out) {

        int max = 0;
        for (SensorReading e : events) {
            if (e.f1 > max) max = e.f1;
        }
        out.collect(new Tuple3<>((Tuple1<String>)key).f0, window.getEnd(), max));
    }
}
```

Window State during Aggregation



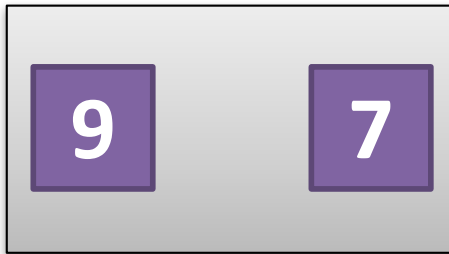
state



Window State during Aggregation



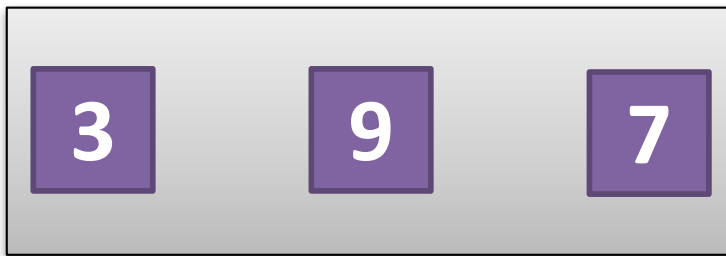
state



Window State during Aggregation




state



Window State during Aggregation



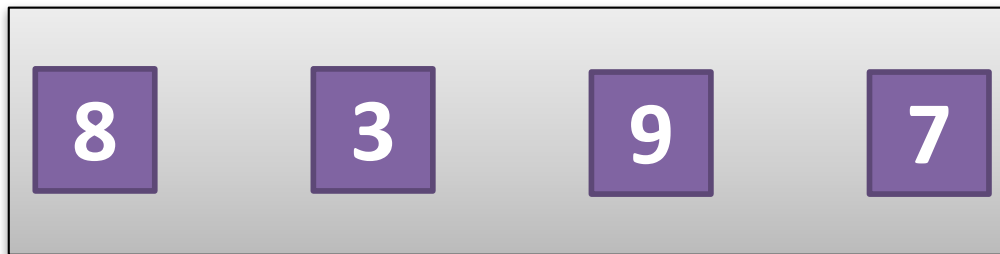
state



Window State during Aggregation



state



window trigger

Incremental Window Aggregation



```
DataStream<SensorReading> input = ...
```

```
input
    .keyBy("key")
    .timeWindow(Time.minutes(1))
    .reduce(new MyReducingMax(), new MyWindowFunction());
```

```
private static class MyReducingMax implements ReduceFunction<SensorReading> {
    public SensorReading reduce(SensorReading r1, SensorReading r2) {
        return r1.value() > r2.value() ? r1 : r2;
    }
}
```

```
private static class MyWindowFunction implements WindowFunction<
    SensorReading, Tuple2<Long, SensorReading>, String, TimeWindow> {
    public void apply(String key,
        TimeWindow window,
        Iterable<SensorReading> maxReadings,
        Collector<Tuple2<Long, SensorReading>> out) {
        SensorReading max= maxReadings.iterator().next();
        out.collect(new Tuple2<Long, SensorReading>(window.getStart(), max));
    }
}
```

Incremental Aggregation



Incremental Aggregation



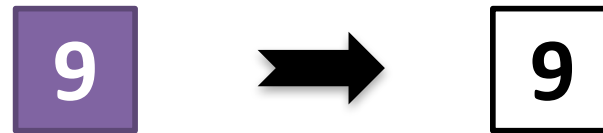
Incremental Aggregation



Incremental Aggregation



Incremental Aggregation



window trigger

Window Operations



- Passed an `Iterable` containing all elements of a `Window`:
 - `apply(windowFunction)`
 - `process(processWindowFunction)`
 - new in 1.3

ProcessWindowFunction()



```
public abstract class ProcessWindowFunction<IN, OUT, KEY, W extends Window> extends AbstractRichFunction {  
  
    /**  
     * Evaluates the window and outputs none or several elements.  
     *  
     * @param key The key for which this window is evaluated.  
     * @param context The context in which the window is being evaluated.  
     * @param elements The elements in the window being evaluated.  
     * @param out A collector for emitting elements.  
     *  
     */  
    public abstract void process(  
        KEY key,  
        Context context,  
        Iterable<IN> elements,  
        Collector<OUT> out) throws Exception;  
  
    // The context holding window metadata.  
    public abstract class Context implements java.io.Serializable {  
        public abstract W window();  
        public abstract long currentProcessingTime();  
        public abstract long currentWatermark();  
        public abstract KeyedStateStore windowState(); // per-key per-window state  
        public abstract KeyedStateStore globalState(); // per-key global state  
    }  
}
```

Incremental Window Operations



- Passed each element of a window, which is aggregated into a single result:
 - `reduce(reduceFunction)`
 - ~~`fold(initialVal, foldFunction)`~~
 - `aggregate(aggregateFunction)`

Other Aggregations



- `sum(key)`, `min(key)`, `max(key)`
 - return the value
- `sumBy(key)`, `minBy(key)`, `maxBy(key)`
 - return an element with the value
- These are available on `KeyedStreams` as well as `WindowedStreams`

Custom window logic

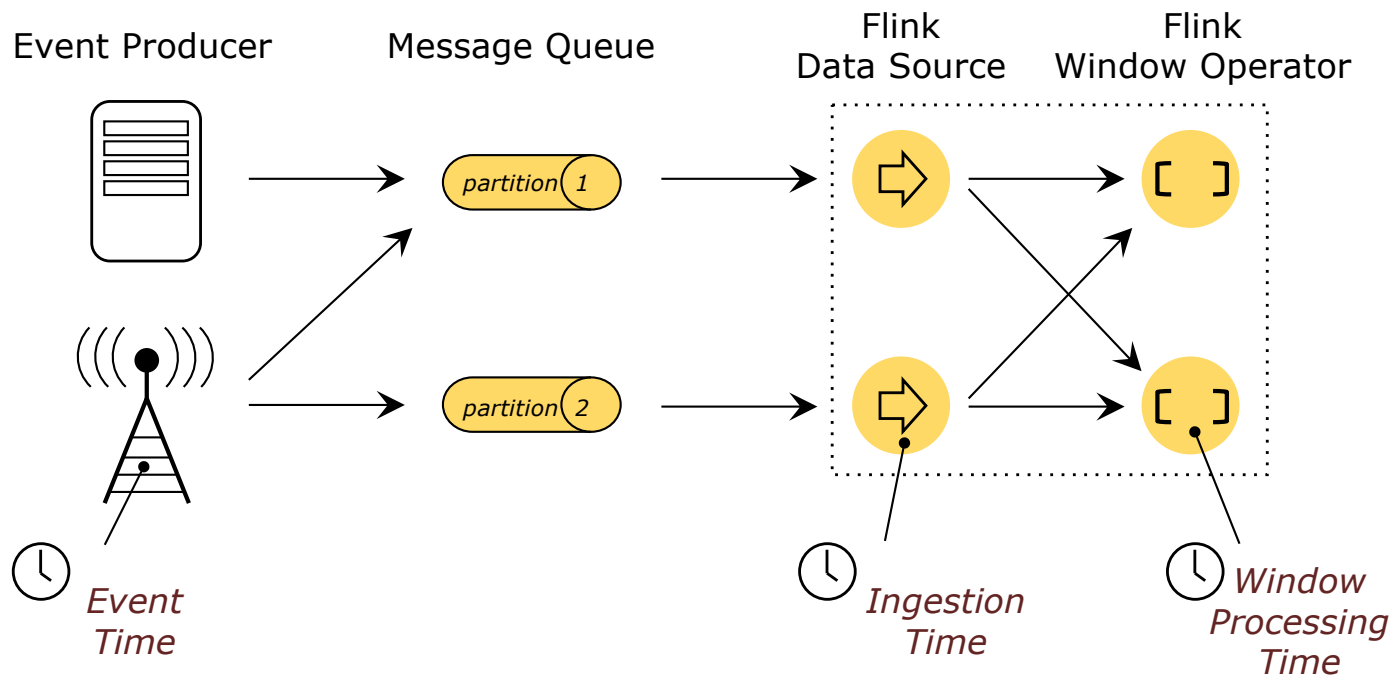


- The DataStream API allows you to define very custom window logic
- **GlobalWindows**
 - a flexible, low-level window assignment scheme that can be used to implement custom windowing behaviors
 - only useful if you explicitly specify triggering, otherwise nothing will happen
- **Trigger**
 - defines when to evaluate a window
 - whether to purge the window or not
- **Careful!** This part of the API requires a good understanding of the windowing mechanism!

Handling Time Explicitly

The **biggest change** in moving from
batch to streaming is
handling time explicitly

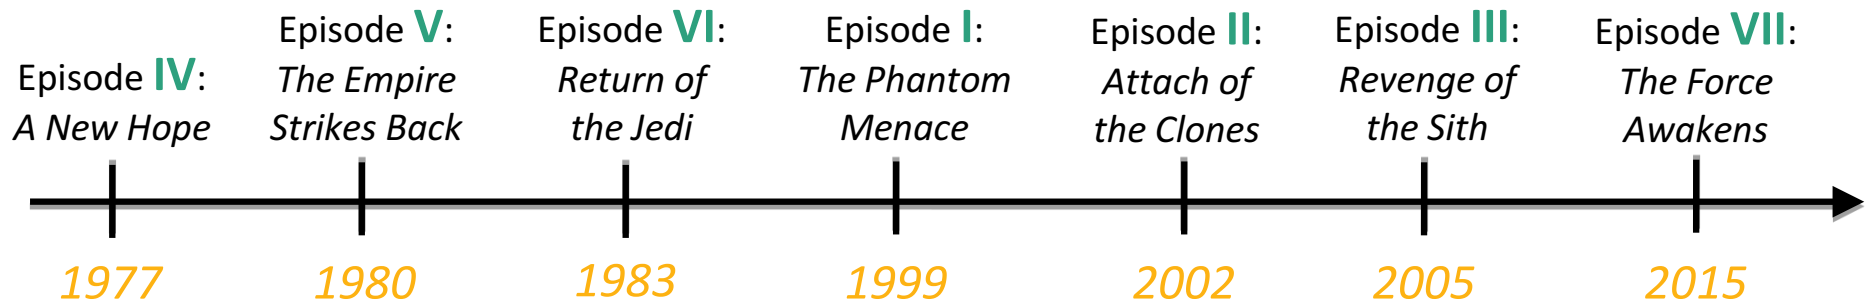
Different Notions of Time



Event Time vs Processing Time



This is called **event time**



This is called **processing time**

Setting the StreamTimeCharacteristic



```
final StreamExecutionEnvironment env =  
    StreamExecutionEnvironment.getExecutionEnvironment();  
  
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);  
  
// alternatively:  
// env.setStreamTimeCharacteristic(TimeCharacteristic.IngestionTime);  
// env.setStreamTimeCharacteristic(TimeCharacteristic.ProcessingTime);
```

Choosing Event Time has Consequences




- With event time, Flink needs to know
 - how to extract timestamps from stream elements
 - when enough event time has elapsed that a time window should be triggered

Watermarks



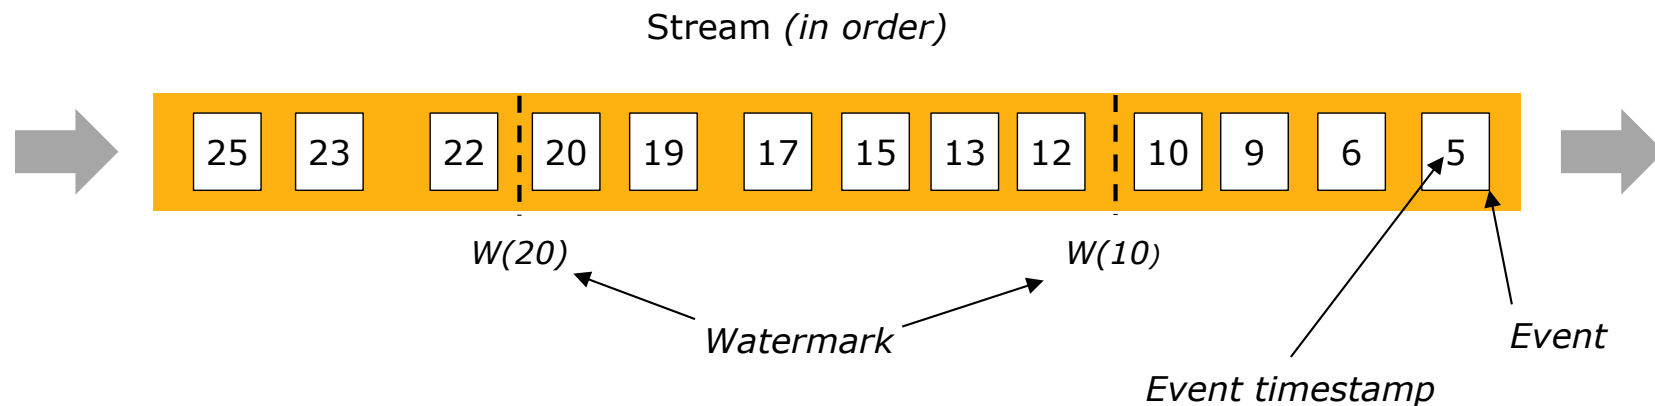
- Watermarks mark the progress of event time
- They flow with the data stream and carry a timestamp
- *Watermarks assert that all earlier events have (probably) arrived*



Perfect Watermarks



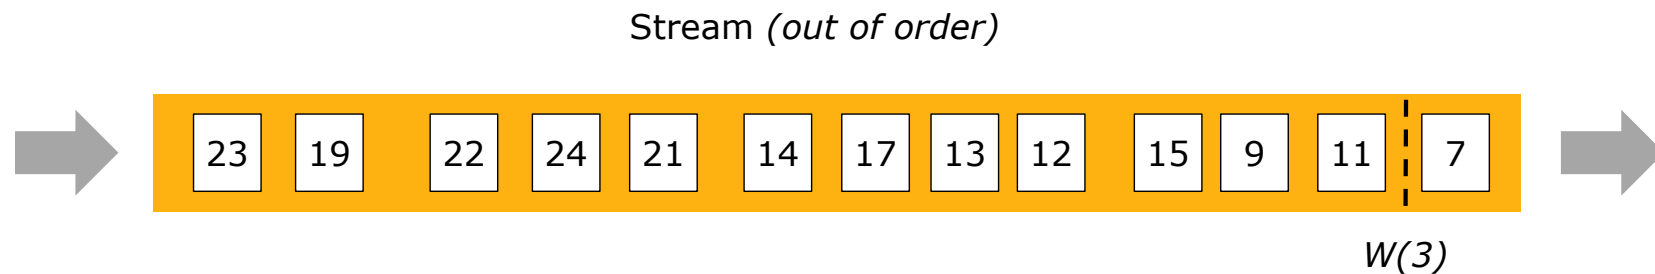
- When stream elements are in order (or in order by key), we can achieve perfect watermarking



Bounded out-of-orderness



- When events are out-of-order, we often assume there is some bound to how out-of-order they can be




$\text{maxOutOfOrderness} = 4$

Bounded out-of-orderness



- Each time a new maximum timestamp arrives, we have enough info to emit a new Watermark

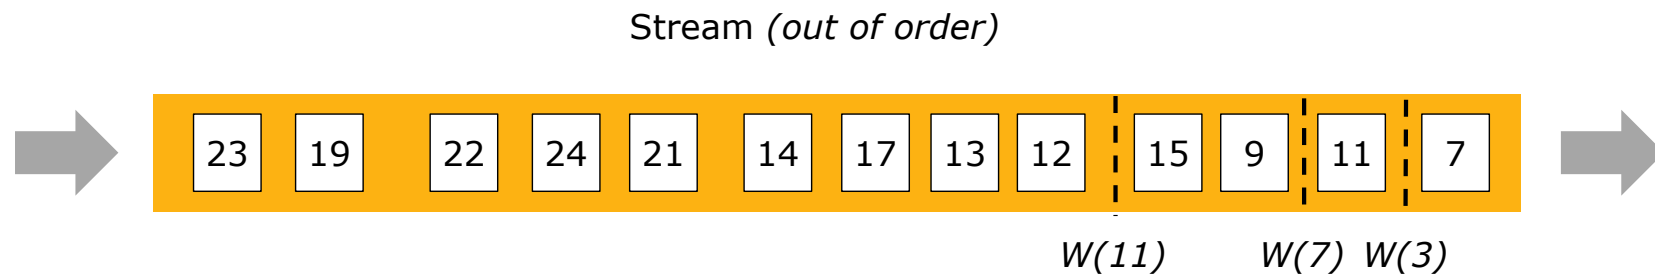


$maxOutOfOrderness = 4$

Bounded out-of-orderness



- Each time a new maximum timestamp arrives, we have enough info to emit a new Watermark




$maxOutOfOrderness = 4$

Bounded out-of-orderness




- Each time a new maximum timestamp arrives, we have enough info to emit a new Watermark



How often to emit Watermarks?




- Here we are emitting a new Watermark as often as possible
- However, it is best to avoid generating too many Watermarks



Watermarks define Lateness



- Elements where $timestamp < currentWatermark$ are late



Two Styles of Watermark Generation



■ Periodic Watermarks

- Based on a timer
- `BoundedOutOfOrdernessGenerator` is an example
- `ExecutionConfig.setAutoWatermarkInterval(msec)` controls the interval at which your periodic watermark generator is called

■ Punctuated Watermarks

- Based on something in the event stream



- **AscendingTimestampExtractor**
 - For special case when timestamps are in ascending order
- **BoundedOutOfOrdernessTimestampExtractor**
 - Periodically emits watermarks that lag a fixed amount of time behind the max timestamp seen so far

Example



```
stream
    .assignTimestampsAndWatermarks(new MyTSExtractor())
    .keyBy(...)
    .timeWindow(...)
    .addSink(...);
```

```
public static class MyTSExtractor extends
    BoundedOutOfOrdernessTimestampExtractor<TaxiRide> {

    public TaxiRideTSExtractor() {
        super(Time.seconds(MAX_EVENT_DELAY));
    }

    @Override
    public long extractTimestamp(TaxiRide ride) {
        return ride.startTime.getMillis();
    }
}
```

```

public class BoundedOutOfOrdernessGenerator extends
    AssignerWithPeriodicWatermarks<MyEvent> {

    private final long maxOutOfOrderness = 3500; // 3.5 seconds

    private long currentMaxTimestamp;

    @Override
    public long extractTimestamp(MyEvent element, long previousElementTimestamp) {
        long timestamp = element.getCreationTime();
        currentMaxTimestamp = Math.max(timestamp, currentMaxTimestamp);
        return timestamp;
    }

    @Override
    public Watermark getCurrentWatermark() {
        // watermark is current highest timestamp minus the out-of-orderness bound
        return new Watermark(currentMaxTimestamp - maxOutOfOrderness);
    }
}

```

```

public class PunctuatedAssigner extends AssignerWithPunctuatedWatermarks<MyEvent> {


    @Override
    public long extractTimestamp(MyEvent element, long previousElementTimestamp) {
        return element.getCreationTime();
    }

    @Override
    public Watermark checkAndGetNextWatermark(MyEvent lastElement,
                                              long extractedTimestamp) {

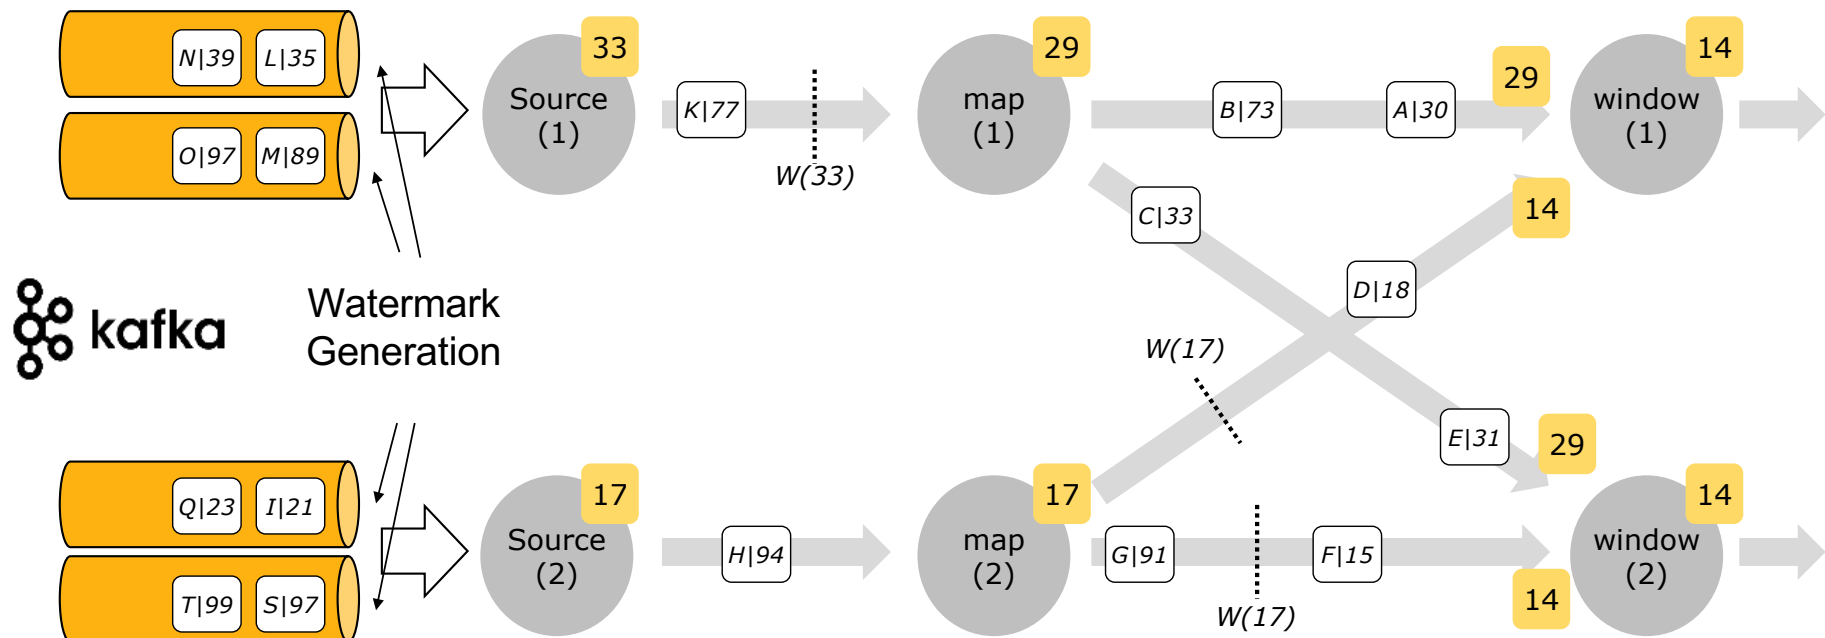
        return lastElement.hasWatermarkMarker() ?
            new Watermark(extractedTimestamp) : null;
    }
}

```

Watermarks in Parallel



Per-Kafka-Partition Watermarks



Watermarking



- Perfect
- (Un)comfortably bounded by fixed delay
 - too slow: results are delayed
 - too fast: some data is late
- Heuristic
 - allow windows to produce results as soon as meaningfully possible, and then continue with updates during the allowed lateness interval